

# CSCE 658: Randomized Algorithms

## Lecture 11

Samson Zhou

# Class Logistics

- **March 5:** Lecture canceled, i.e., do NOT show up to HRBB 126 (unless you want to see an empty classroom)

# Previously in the Streaming Model

- Reservoir sampling
- Heavy-hitters
  - Misra-Gries
  - CountMin
  - CountSketch
- Moment estimation
  - AMS algorithm
- Sparse recovery
- Distinct elements estimation

# Reservoir Sampling

- Suppose we see a stream of elements from  $[n]$ . How do we uniformly sample one of the positions of the stream?

47 72 81 10 14 33 51 29 54 9 36 46 10

# Heavy-Hitters (Frequent Items)

- Given a set  $S$  of  $m$  elements from  $[n]$ , let  $f_i$  be the frequency of element  $i$ . (How often it appears)
- Let  $L_p$  be the norm of the frequency vector:

$$L_p = (f_1^p + f_2^p + \dots + f_n^p)^{1/p}$$

- **Goal:** Given a set  $S$  of  $m$  elements from  $[n]$  and a threshold  $\varepsilon$ , output the elements  $i$  such that  $f_i > \varepsilon L_p$ ...and no elements  $j$  such that  $f_j < \frac{\varepsilon}{2} L_p$  (we saw algorithms for  $p = 1$  and  $p = 2$ )
- **Motivation:** DDoS prevention, iceberg queries

# Frequency Moments ( $L_p$ Norm)

- Given a set  $S$  of  $m$  elements from  $[n]$ , let  $f_i$  be the frequency of element  $i$ . (How often it appears)
- Let  $F_p$  be the frequency moment of the vector:

$$F_p = f_1^p + f_2^p + \cdots + f_n^p$$

- **Goal:** Given a set  $S$  of  $m$  elements from  $[n]$  and an accuracy parameter  $\varepsilon$ , output a  $(1 + \varepsilon)$ -approximation to  $F_p$
- **Motivation:** Entropy estimation, linear regression

# The Streaming Model


- So far, all questions have been *statistical*
- What other questions can be asked? (Think in general, outside of the streaming model)

# The Streaming Model

- So far, all questions have been *statistical*
- What other questions can be asked? (Think in general, outside of the streaming model)
- Algebraic, geometric



# The Streaming Model

- So far, all questions have been *statistical*
  - What other questions can be asked? (Think in general, outside of the streaming model)
  - Algebraic, geometric
- TODAY
- 

# Graph Theory


- Suppose we have a graph  $G$  with vertex set  $V$  and edge set  $E$
- Let  $V = [n]$  for simplicity, so each vertex is an integer from  $1$  to  $n$
- Then each edge  $e \in E$  can be written as  $e = (u, v)$  for  $u, v \in [n]$
- In other words, each edge is a pair of integers from  $1$  to  $n$

# Graph Theory

- For today, we will assume a simple, undirected, unweighted graph
- Graph has no self-loops, no multi-edges
- Edges are undirected
- Each edge has weight **1**


# Semi-streaming Model


- Recall that we have a graph  $G = (V = [n], E)$
- Suppose  $|E| = m$
- The edges of the graph arrive sequentially, i.e., insertion-only model
- We are allowed to use  $n \cdot \text{polylog}(n)$  space
- Enough to store things like a matching, a spanning tree, **NOT** enough to store entire graph, since  $m$  can be as large as  $O(n^2)$




# Bipartiteness


- **Bipartite graph**: Graph can be partitioned into two disjoint sets  $L$  and  $R$  so that every edge is between a vertex in  $L$  and a vertex in  $R$
- **Goal**: Given a graph  $G$ , determine whether  $G$  is a bipartite graph













# Applications for Bipartiteness Testing

- **Graph coloring:** You want to color a graph such that no neighboring items share the same color



# Applications for Bipartiteness Testing

- **Circuit design:** In electrical engineering and VLSI (Very Large Scale Integration) design, you may want to know if a circuit can be optimally partitioned into two complementary parts, which can be achieved by testing the bipartiteness of the circuit's dependency graph



# Bipartiteness

- What is a necessary and sufficient condition for bipartiteness?

# Bipartiteness

- What is a necessary and sufficient condition for bipartiteness?
- A graph is bipartite if and only if it can be colored using two colors (a coloring of a graph is an assignment of colors to vertices such that no two vertices share the same color)
- A graph is bipartite if and only if it has no odd cycles


# Bipartiteness


- How to perform bipartiteness testing in the central setting?


# Bipartiteness


- How to perform bipartiteness testing in the central setting?
- Start at arbitrary vertex, run BFS, and assign alternating levels to different side until there is a contradiction











# Bipartiteness in the Streaming Model

- Bipartiteness is a monotone property, i.e., additional edges to a graph that is not bipartite will result in a graph that is not bipartite

# Bipartiteness in the Streaming Model

- **Intuition:** Greedily add edges to minimum spanning forest
- **Algorithm:**
  1. Initialize  $F = \emptyset$ .
  2. For each edge  $e = (u, v)$ :
    1. If  $F \cup (u, v)$  does not contain a cycle, add  $(u, v)$  to  $F$ :  $F \leftarrow F \cup (u, v)$
    2. If  $F \cup (u, v)$  contains an odd cycle, return GRAPH IS NOT BIPARTITE
  3. Return GRAPH IS BIPARTITE




# Bipartiteness in the Streaming Model

- Algorithm maintains a tree (because it does not add any edges that would create cycles)
- How many edges does the algorithm keep?

# Bipartiteness in the Streaming Model


- Algorithm maintains a tree (because it does not add any edges that would create cycles)
- Algorithm can keep at most  $n$  edges, so the total space usage is  $O(n)$  words of space.






# Clustering

- **Goal:** Given input dataset  $X$ , partition  $X$  so that “similar” points are in the same cluster and “different” points are in different clusters




# $k$ -Clustering

- **Goal:** Given input dataset  $X$ , partition  $X$  so that “similar” points are in the same cluster and “different” points are in different clusters
- There can be at most  $k$  different clusters



# $k$ -Clustering

- **Question:** How do we measure the “quality” of each clustering?



# $k$ -Clustering

- **Question:** How do we measure the “quality” of each clustering?
- Assign a “center”  $c_i$  to each cluster
- Have a cost function induced by  $c_i$  for all of the points  $P_i$  assigned to cluster  $i$

# $k$ -Clustering


- **Question**: How do we measure the “quality” of each clustering?
- Assign a “center”  $c_i$  to each cluster
- Have a cost function induced by  $c_i$  for all of the points  $P_i$  assigned to cluster  $i$ 
  - Assume points are in metric space with distance function  $\text{dist}(\cdot, \cdot)$
  - Define  $\text{Cost}(P_i, c_i)$  to be a function of  $\{\text{dist}(x, c_i)\}_{x \in P_i}$

# $k$ -Clustering

- **Question:** How do we measure the “quality” of each clustering?
- Have a cost function induced by  $c_i$  for all of the points  $P_i$  assigned to cluster  $i$ 
  - Define  $\text{Cost}(P_i, c_i)$  to be a function of  $\{\text{dist}(x, c_i)\}_{x \in P_i}$
- Suppose the set of centers is  $C = \{c_1, \dots, c_k\}$ 
  - Define clustering cost  $\text{Cost}(X, C)$  to be a function of  $\{\text{dist}(x, C)\}_{x \in X}$

# $k$ -Clustering


- Define clustering cost  $\text{Cost}(X, C)$  to be a function of  $\{\text{dist}(x, C)\}_{x \in C}$






# $k$ -Clustering

- Define clustering cost  $\text{Cost}(X, C)$  to be a function of  $\{\text{dist}(x, C)\}_{x \in X}$
- $k$ -center:  $\text{Cost}(X, C) = \max_{x \in X} \text{dist}(x, C)$




# $k$ -Clustering

- Define clustering cost  $\text{Cost}(X, C)$  to be a function of  $\{\text{dist}(x, C)\}_{x \in X}$
- $k$ -center:  $\text{Cost}(X, C) = \max_{x \in X} \text{dist}(x, C)$
- $k$ -median:  $\text{Cost}(X, C) = \sum_{x \in X} \text{dist}(x, C)$




# $k$ -Clustering

- Define clustering cost  $\text{Cost}(X, C)$  to be a function of  $\{\text{dist}(x, C)\}_{x \in X}$
- $k$ -center:  $\text{Cost}(X, C) = \max_{x \in X} \text{dist}(x, C)$
- $k$ -median:  $\text{Cost}(X, C) = \sum_{x \in X} \text{dist}(x, C)$
- $k$ -means:  $\text{Cost}(X, C) = \sum_{x \in X} (\text{dist}(x, C))^2$



# $k$ -Clustering

- Define clustering cost  $\text{Cost}(X, C)$  to be a function of  $\{\text{dist}(x, C)\}_{x \in X}$
- $k$ -center:  $\text{Cost}(X, C) = \max_{x \in X} \text{dist}(x, C)$
- $k$ -median:  $\text{Cost}(X, C) = \sum_{x \in X} \text{dist}(x, C)$
- $k$ -means:  $\text{Cost}(X, C) = \sum_{x \in X} (\text{dist}(x, C))^2$
- $(k, z)$ -clustering:  $\text{Cost}(X, C) = \sum_{x \in X} (\text{dist}(x, C))^z$



# Euclidean $k$ -Clustering

- For Euclidean  $k$ -clustering, input points  $X = x_1, \dots, x_n$  are in  $\mathbb{R}^d$  (for us, they will be in  $[\Delta]^d := \{1, 2, \dots, \Delta\}^d$ )
- $\text{dist}(x, y) = \sqrt{(x_1 - y_1)^2 + \dots + (x_d - y_d)^2}$  is the Euclidean distance
- $(k, z)$ -clustering problem:

$$\min_{C: |C| \leq k} \text{Cost}(X, C) = \min_{C: |C| \leq k} \sum_{x \in X} (\text{dist}(x, C))^z$$

$(-8, 4)$



$(-4, 3)$



$(4, 3)$

$(-3, 0)$

$(-11, -4)$



$(0, -4)$










$k$ -center:  $\text{Cost}(X, C) = \max_{x \in X} \text{dist}(x, C) = 5$



$k$ -median:  $\text{Cost}(X, C) = \sum_{x \in X} \text{dist}(x, C) = 4 + 5 + 5 + 3 + 4 + 5 = 26$



$k$ -means:  $\text{Cost}(X, C) = \sum_{x \in X} (\text{dist}(x, C))^2 = 16 + 25 + 25 + 9 + 16 + 25 = 116$






# Coreset

- Subset  $X'$  of representative points of  $X$  for a specific clustering objective
- $\text{Cost}(X, C) \approx \text{Cost}(X', C)$   
for all sets  $C$  with  $|C| = k$




# Coreset

- Subset  $X'$  of representative points of  $X$  for a specific clustering objective
- $\text{Cost}(X, C) \approx \text{Cost}(X', C)$   
for all sets  $C$  with  $|C| = k$



# Coreset

- Subset  $X'$  of representative points of  $X$  for a specific clustering objective
- $\text{Cost}(X, C) \approx \text{Cost}(X', C)$   
for all sets  $C$  with  $|C| = k$



# Coreset (Formal Definition)

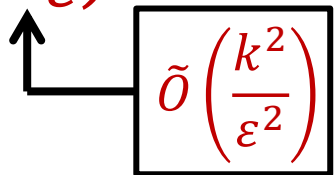
- Given a set  $X$  and an accuracy parameter  $\varepsilon > 0$ , we say a set  $X'$  with weight function  $w$  is an  $(1 + \varepsilon)$ -*multiplicative coreset* for a cost function  $\text{Cost}$ , if for all queries  $C$  with  $|C| \leq k$ , we have

$$(1 - \varepsilon)\text{Cost}(X, C) \leq \text{Cost}(X', C, w) \leq (1 + \varepsilon)\text{Cost}(X, C)$$



$$(k, z)\text{-clustering: } \text{Cost}(X', C, w) = \sum_{x \in X'} w(x) \cdot (\text{dist}(x, C))^z$$

# $(k, z)$ -Clustering in the Streaming Model

- Merge-and-reduce framework
- Suppose there exists a  $(1 + \varepsilon)$ -coreset construction for  $(k, z)$ -clustering that uses  $f\left(k, \frac{1}{\varepsilon}\right)$  weighted input points  

- Partition the stream into blocks containing  $f\left(k, \frac{\log n}{\varepsilon}\right)$  points



# $(k, z)$ -Clustering in the Streaming Model

- Partition the stream into blocks containing  $f\left(k, \frac{\log n}{\varepsilon}\right)$  points
- Create a  $\left(1 + \frac{\varepsilon}{\log n}\right)$ -coreset for each block
- Create a  $\left(1 + \frac{\varepsilon}{\log n}\right)$ -coreset for the set of points formed by the union of two coresets for each block

Reduce




Merge




# $(k, z)$ -Clustering in the Streaming Model

- Partition the stream into blocks containing  $f\left(k, \frac{\log n}{\varepsilon}\right)$  points
- Create a  $\left(1 + \frac{\varepsilon}{\log n}\right)$ -coreset for each block
- Create a  $\left(1 + \frac{\varepsilon}{\log n}\right)$ -coreset for the set of points formed by the union of two coresets for each block



# $(k, z)$ -Clustering in the Streaming Model

- There are  $O(\log n)$  levels
- Each coreset is a  $\left(1 + \frac{\varepsilon}{\log n}\right)$ -coreset of two coresets
- Total approximation is  $\left(1 + \frac{\varepsilon}{\log n}\right)^{\log n} = (1 + O(\varepsilon))$



# $(k, z)$ -Clustering in the Streaming Model

- Suppose there exists a  $(1 + \varepsilon)$ -coreset construction for  $(k, z)$ -clustering that uses  $f\left(k, \frac{1}{\varepsilon}\right)$  weighted input points
- Partition the stream into blocks containing  $f\left(k, \frac{\log n}{\varepsilon}\right)$  points
- Total space is  $f\left(k, \frac{\log n}{\varepsilon}\right) \cdot O(\log n)$  points

For  $k$ -means clustering, this is  $\tilde{O}\left(\frac{k^2}{\varepsilon^2} \cdot \log^3 n\right)$  points