


Tight Bounds for Adversarially Robust Streams and Sliding Windows via Difference Estimators




Samson Zhou

**Carnegie
Mellon
University**

(Based on work joint with David P.
Woodruff)

Background

- ❖ Post-doc at Carnegie Mellon
- ❖ PhD in Computer Science from Purdue
- ❖ Bachelors in Math, Computer Science from MIT



- ❖ **Research areas:** Security and Privacy, Data Science, Sublinear Algorithms

Recent Work on Security and Privacy

- Private Data Stream Analysis for Universal Symmetric Norm Estimation (FORC 2022)
- On the Security of Proofs of Sequential Work in a Post-Quantum World (ITC 2021)
- Computationally Data-Independent Memory Hard Functions (ITCS 2020)
- Data-Independent Memory Hard Functions: New Attacks and Stronger Constructions (CRYPTO 2019)
- Bandwidth-Hard Functions: Reductions and Lower Bounds (CCS 2018)
- On the Economics of Offline Password Cracking (Security and Privacy 2018)

Recent Work on Data Science

- Learning-Augmented k-means Clustering (ICLR 2022)
- Fast Regression for Structured Inputs (ICLR 2022)
- New Coresets for Projective Clustering and Applications (AISTATS 2022)
- Dimensionality Reduction for Wasserstein Barycenter (NeurIPS 2021)
- Learning a Latent Simplex in Input Sparsity Time (ICLR 2021)
- Near Optimal Linear Algebra in the Online and Sliding Window Models (FOCS 2020)
- Data-Independent Neural Pruning via Coresets (ICLR 2020)
- Adversarially Robust Submodular Maximization under Knapsack Constraints (KDD 2019)

Recent Work on Sublinear Algorithms

- Memory Bounds for the Experts Problem (STOC 2022)
- The White-Box Adversarial Data Stream Model (PODS 2022)
- Truly Perfect Samplers for Data Streams and Sliding Windows (PODS 2022)
- Noisy Boolean Hidden Matching with Applications (ITCS 2022)
- Adversarial Robustness of Streaming Algorithms through Importance Sampling (NeurIPS 2021)
- Tight Bounds for Adversarially Robust Streams and Sliding Windows via Difference Estimators (FOCS 2021)
- Separations for Estimating Large Frequency Moments on Data Streams (ICALP 2021)
- Non-Adaptive Adaptive Sampling on Turnstile Streams (STOC 2020)

Model #1: Streaming Model

- ❖ **Input:** Elements of an underlying data set S , which arrives sequentially
- ❖ **Output:** Evaluation (or approximation) of a given function
- ❖ **Goal:** Use space *sublinear* in the size m of the input S

1 0 1 1 1 0 0 1

Heavy-Hitters

- ❖ Given a set S of m elements from $[n]$, let f_i be the frequency of element i . (How often it appears)

$$1\ 1\ 2\ 1\ 2\ 1\ 1\ 2\ 3 \rightarrow [5, 3, 1, 0] := f$$

Heavy-Hitters

- ❖ Given a set S of m elements from $[n]$, let f_i be the frequency of element i . (How often it appears)
- ❖ Let L_2 be the norm of the frequency vector:

$$L_2 = \sqrt{f_1^2 + f_2^2 + \dots + f_n^2}$$

- ❖ **Goal**: Given a set S of m elements from $[n]$ and a threshold ϵ , output the elements i such that $f_i > \epsilon L_2$...and no elements j such that $f_j < \frac{\epsilon}{16} L_2$
- ❖ **Motivation**: DDoS prevention, iceberg queries

Frequency Moments

- ❖ Given a set S of m elements from $[n]$, let f_i be the frequency of element i . (How often it appears)
- ❖ Let F_p be the frequency moment of the vector:

$$F_p = f_1^p + f_2^p + \cdots + f_n^p$$

- ❖ **Goal**: Given a set S of m elements from $[n]$ and an accuracy parameter ϵ , output a $(1 + \epsilon)$ -approximation to F_p
- ❖ **Motivation**: Entropy estimation, linear regression

Distinct Elements

- ❖ Given a set S of m elements from $[n]$, let f_i be the frequency of element i . (How often it appears)
- ❖ Let F_0 be the frequency moment of the vector:

$$F_0 = |\{i : f_i \neq 0\}|$$

- ❖ **Goal**: Given a set S of m elements from $[n]$ and an accuracy parameter ϵ , output a $(1 + \epsilon)$ -approximation to F_0
- ❖ **Motivation**: Traffic monitoring

$(1 + \epsilon)$ -Approximation Streaming Algorithms

- ❖ Space $O\left(\frac{1}{\epsilon^2} + \log n\right)$ algorithm for F_0 [KaneNelsonWoodruff10], [Blasiok20]
- ❖ Space $O\left(\frac{1}{\epsilon^2} \log n\right)$ algorithm for F_p with $p \in (0, 2]$ [BlasiokDingNelson17]
- ❖ Space $O\left(\frac{1}{\epsilon^2} n^{1-2/p} \log^2 n\right)$ algorithm for F_p with $p > 2$ [Ganguly11, GangulyWoodruff18]
- ❖ Space $O\left(\frac{1}{\epsilon^2} \log n\right)$ algorithm for L_2 -heavy hitters [BravermanChestnutIvkinNelsonWangWoodruff17]

Model #2: Adversarially Robust Streaming

- ❖ **Input:** Elements of an underlying data set S , which arrives sequentially and *adversarially*
- ❖ **Output:** Evaluation (or approximation) of a given function
- ❖ **Goal:** Use space *sublinear* in the size m of the input S



1

1



Model #2: Adversarially Robust Streaming

- ❖ **Input:** Elements of an underlying data set S , which arrives sequentially and *adversarially*
- ❖ **Output:** Evaluation (or approximation) of a given function
- ❖ **Goal:** Use space *sublinear* in the size m of the input S



10

1



Model #2: Adversarially Robust Streaming

- ❖ **Input:** Elements of an underlying data set S , which arrives sequentially and *adversarially*
- ❖ **Output:** Evaluation (or approximation) of a given function
- ❖ **Goal:** Use space *sublinear* in the size m of the input S



101

2



Model #2: Adversarially Robust Streaming

- ❖ **Input:** Elements of an underlying data set S , which arrives sequentially and *adversarially*
- ❖ **Output:** Evaluation (or approximation) of a given function
- ❖ **Goal:** Use space *sublinear* in the size m of the input S



1010

3



Model #2: Adversarially Robust Streaming

- ❖ **Input:** Elements of an underlying data set S , which arrives sequentially and *adversarially*
- ❖ **Output:** Evaluation (or approximation) of a given function
- ❖ **Goal:** Use space *sublinear* in the size m of the input S



1010

3



Model #2: Adversarially Robust Streaming

- ❖ **Input:** Elements of an underlying data set S , which arrives sequentially and *adversarially*
- ❖ **Output:** Evaluation (or approximation) of a given function
- ❖ **Goal:** Use space *sublinear* in the size m of the input S

- ❖ **Adversarially Robust:** “Future queries may depend on previous queries”
- ❖ **Motivation:** Database queries, adversarial ML

$(1 + \epsilon)$ -Robust Algorithms [Ben-EliezerJayaramWoodruffYogev20]

- ❖ Space $\tilde{O}\left(\frac{1}{\epsilon^3} \log n\right)$ algorithm for F_0
- ❖ Space $\tilde{O}\left(\frac{1}{\epsilon^3} \log n\right)$ algorithm for F_p with $p \in (0, 2]$
- ❖ Space $\tilde{O}\left(\frac{1}{\epsilon^3} n^{1-2/p}\right)$ algorithm for F_p with $p > 2$
- ❖ Space $\tilde{O}\left(\frac{1}{\epsilon^3} \log n\right)$ algorithm for L_2 -heavy hitters

“A general framework that loses* nothing in n and only $\frac{1}{\epsilon}$ ”

“What’s an epsilon between friends?”

- ❖ Statista: $\sim 300B \sim 2^{38}$ e-mails sent per day
- ❖ Accuracy: $\epsilon = 0.01$
- ❖ Since $\frac{1}{\epsilon} > \log n$, we should care about $\frac{1}{\epsilon}$ factors!

$(1 + \epsilon)$ -Robust Algorithms

[HassidimKaplanMansourMatiasStemmer20]

- ❖ Space $\tilde{O}\left(\frac{1}{\epsilon^{2.5}} \log^4 n\right)$ algorithm for F_0
- ❖ Space $\tilde{O}\left(\frac{1}{\epsilon^{2.5}} \log^4 n\right)$ algorithm for F_p with $p \in (0, 2]$
- ❖ Space $\tilde{O}\left(\frac{1}{\epsilon^{2.5}} n^{1-2/p}\right)$ algorithm for F_p with $p > 2$
- ❖ Space $\tilde{O}\left(\frac{1}{\epsilon^{2.5}} \log^4 n\right)$ algorithm for L_2 -heavy hitters

“ $\frac{1}{\epsilon}$ losses are not necessary”

Our Results: $(1 + \epsilon)$ -Robust Algorithms

- ❖ Space $\tilde{O}\left(\frac{1}{\epsilon^2} \log n\right)$ algorithm for F_0
- ❖ Space $\tilde{O}\left(\frac{1}{\epsilon^2} \log n\right)$ algorithm for F_p with $p \in (0, 2]$
- ❖ Space $\tilde{O}\left(\frac{1}{\epsilon^2} n^{1-2/p}\right)$ algorithm for F_p with integer $p > 2$
- ❖ Space $\tilde{O}\left(\frac{1}{\epsilon^2} \log n\right)$ algorithm for L_2 -heavy hitters

“No losses* are necessary!”

Summary: $(1 + \epsilon)$ -Robust Algorithms

| Problem | [BJWY20] Space | [HKM ⁺ 20] Space | Our Result |
|--|---|---|---|
| Distinct Elements | $\tilde{O}\left(\frac{\log n}{\epsilon^3}\right)$ | $\tilde{O}\left(\frac{\log^4 n}{\epsilon^{2.5}}\right)$ | $\tilde{O}\left(\frac{\log n}{\epsilon^2}\right)$ |
| F_p Estimation, $p \in (0, 2]$ | $\tilde{O}\left(\frac{\log n}{\epsilon^3}\right)$ | $\tilde{O}\left(\frac{\log^4 n}{\epsilon^{2.5}}\right)$ | $\tilde{O}\left(\frac{\log n}{\epsilon^2}\right)$ |
| Shannon Entropy | $\tilde{O}\left(\frac{\log^6 n}{\epsilon^5}\right)$ | $\tilde{O}\left(\frac{\log^4 n}{\epsilon^{3.5}}\right)$ | $\tilde{O}\left(\frac{\log^3 n}{\epsilon^2}\right)$ |
| L_2 -Heavy Hitters | $\tilde{O}\left(\frac{\log n}{\epsilon^3}\right)$ | $\tilde{O}\left(\frac{\log^4 n}{\epsilon^{2.5}}\right)$ | $\tilde{O}\left(\frac{\log n}{\epsilon^2}\right)$ |
| F_p Estimation, integer $p > 2$ | $\tilde{O}\left(\frac{n^{1-2/p}}{\epsilon^3}\right)$ | $\tilde{O}\left(\frac{n^{1-2/p}}{\epsilon^{2.5}}\right)$ | $\tilde{O}\left(\frac{n^{1-2/p}}{\epsilon^2}\right)$ |
| F_p Estimation, $p \in (0, 2]$, flip number λ | $\tilde{O}\left(\frac{\lambda \log^2 n}{\epsilon^2}\right)$ | $\tilde{O}\left(\frac{\log^3 n \sqrt{\lambda \log n}}{\epsilon^2}\right)$ | $\tilde{O}\left(\frac{\lambda \log^2 n}{\epsilon}\right)$ |

Model #3: Sliding Window Model

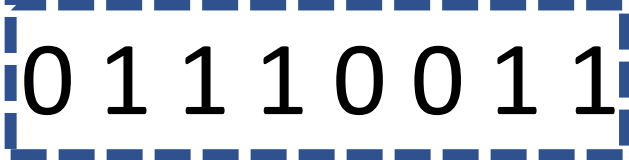
- ❖ **Input:** Elements of an underlying data set S , which arrives sequentially
- ❖ **Output:** Evaluation (or approximation) of a given function
- ❖ **Goal:** Use space *sublinear* in the size m of the input S
- ❖ **Sliding Window:** “Only the m most recent updates form the underlying data set S ”

1 0 1 1 1 0 0 1

Model #3: Sliding Window Model

- ❖ **Input:** Elements of an underlying data set S , which arrives sequentially
- ❖ **Output:** Evaluation (or approximation) of a given function
- ❖ **Goal:** Use space *sublinear* in the size m of the input S
- ❖ **Sliding Window:** “Only the m most recent updates form the underlying data set S ”

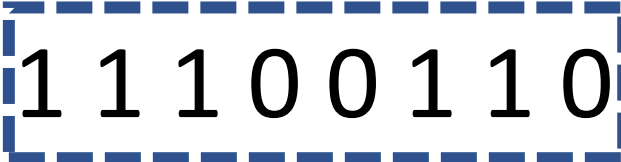
1 0 1 1 1 0 0 1 1



Model #3: Sliding Window Model

- ❖ **Input:** Elements of an underlying data set S , which arrives sequentially
- ❖ **Output:** Evaluation (or approximation) of a given function
- ❖ **Goal:** Use space *sublinear* in the size m of the input S
- ❖ **Sliding Window:** “Only the m most recent updates form the underlying data set S ”


1 0 1 1 1 0 0 1 1 0



Model #3: Sliding Window Model

- ❖ **Input:** Elements of an underlying data set S , which arrives sequentially
- ❖ **Output:** Evaluation (or approximation) of a given function
- ❖ **Goal:** Use space *sublinear* in the size m of the input S
- ❖ **Sliding Window:** “Only the m most recent updates form the underlying data set S ”
 - ❖ Emphasizes recent interactions, appropriate for time sensitive settings

1 0 1 1 1 0 0 1 1 0 1



$(1 + \epsilon)$ -Approximation Sliding Window Algorithms

- ❖ Space $\tilde{O}\left(\frac{1}{\epsilon^2} \log n\right)$ algorithm for F_0
[BravermanGrigorescuLangWoodruffZhou18]
- ❖ Space $\tilde{O}\left(\frac{1}{\epsilon^2} \log^3 n\right)$ algorithm for L_2 -heavy hitters
[BravermanGrigorescuLangWoodruffZhou18]

$(1 + \epsilon)$ -Approximation Sliding Window Algorithms

- ❖ Space $O\left(\frac{1}{\epsilon^3} \log^3 n\right)$ algorithm for F_p with $p \in (0,1)$
[BravermanOstrovsky07]
- ❖ Space $O\left(\frac{1}{\epsilon^{2+p}} \log^3 n\right)$ algorithm for F_p with $p \in (1,2]$
[BravermanOstrovsky07]
- ❖ Space $\tilde{O}\left(\frac{1}{\epsilon^{2+p}} n^{1-2/p}\right)$ algorithm for F_p with $p > 2$
[BravermanOstrovsky07]

“A general framework that loses* nothing in n and only $\frac{1}{\epsilon}$ ”

Our Results: $(1 + \epsilon)$ -Approximation Sliding Window Algorithms

❖ Space $\tilde{O}\left(\frac{1}{\epsilon^2} \log^3 n\right)$ algorithm for F_p with $p \in (0, 2]$

| Problem | [BO07] Space | Our Result |
|-----------------------------------|--|--|
| L_p Estimation, $p \in (0, 1)$ | $\tilde{O}\left(\frac{\log^3 n}{\epsilon^3}\right)$ | $\tilde{O}\left(\frac{\log^3 n}{\epsilon^2}\right)$ |
| L_p Estimation, $p \in (1, 2]$ | $\tilde{O}\left(\frac{\log^3 n}{\epsilon^{2+p}}\right)$ | $\tilde{O}\left(\frac{\log^3 n}{\epsilon^2}\right)$ |
| L_p Estimation, integer $p > 2$ | $\tilde{O}\left(\frac{n^{1-2/p}}{\epsilon^{2+p}}\right)$ | $\tilde{O}\left(\frac{n^{1-2/p}}{\epsilon^2}\right)$ |
| Entropy Estimation | $\tilde{O}\left(\frac{\log^5 n}{\epsilon^4}\right)$ | $\tilde{O}\left(\frac{\log^5 n}{\epsilon^2}\right)$ |

“ $\frac{1}{\epsilon}$ losses are not necessary”

Format

- ❖ Part 1: Background
- ❖ Part 2: Frameworks
- ❖ Part 3: Difference Estimators

Questions?



AMS F_2 Algorithm

- ❖ Let $s \in \{-1, +1\}^n$ be a sign vector of length n
- ❖ Let $Z = \langle s, f \rangle = s_1 f_1 + \dots + s_n f_n$ and consider Z^2

$$E[Z^2] = \sum_{i,j} E[s_i s_j f_i f_j] = f_1^2 + \dots + f_n^2$$

$$\text{Var}[Z^2] \leq \sum_{i,j} E[s_i s_j s_k s_l f_i f_j f_k f_l] \leq 2F_2^2$$

- ❖ Take the mean of $O\left(\frac{1}{\epsilon^2}\right)$ inner products for $(1 + \epsilon)$ -approximation
[AlonMatiasSzegedy99]

“Attack” on AMS

- ❖ Can learn whether $s_i = s_j$ from $\langle s, e_i + e_j \rangle$
- ❖ Let $f_i = 1$ if $s_i = s_1$ and $f_i = -1$ if $s_i \neq s_1$
- ❖ $Z = \langle s, f \rangle = s_1 f_1 + \dots + s_n f_n = m$ and $Z^2 = m^2$ deterministically
- ❖ What happened? Randomness of algorithm not independent of input

Reconstruction Attack on Linear Sketches

- ❖ Linear sketches are “not robust” to adversarial attacks, must use $\Omega(n)$ space [HardtWoodruff13]
- ❖ Approximately learn sketch matrix U , query something in the kernel of U
- ❖ Iterative process, start with V_1, \dots, V_r
- ❖ **Correlation finding**: Find vectors weakly correlated with U orthogonal to V_{i-1}
- ❖ **Boosting**: Use these vectors to find strongly correlated vector v
- ❖ **Progress**: Set $V_i = \text{span}(V_{i-1}, v)$

Insertion-Only Streams


- ❖ **Key:** Deletions are needed to perform this attack
- ❖ Similar lower bounds for the sliding window model
[\[DatarGionisIndykMotwani02\]](#)
- ❖ Assume insertion-only updates
- ❖ How do the previous results work?

Robust Algorithms

- ❖ Suppose we are trying to approximate some given function
 1. Suppose we have a streaming algorithm for this function
 2. Suppose this function is monotonic and the stream is insertion-only
- ❖ Sketch switching framework [\[Ben-EliezerJayaramWoodruffYogev20\]](#) gives a robust for this function
- ❖ Start many instances of the streaming algorithm at the beginning
- ❖ Use an instance of the algorithm but “freeze” the output
- ❖ Each time the next instance has value $(1 + O(\epsilon))$ more than the “frozen” output, use the next instance and “freeze” its output

Robust Algorithms


- ❖ Start many instances of the streaming algorithm at the beginning
- ❖ Use an instance of the algorithm but “freeze” the output
- ❖ Each time the next instance has value $(1 + O(\epsilon))$ more than the “frozen” output, use the next instance and “freeze” its output



- ❖ Example: Number of ones in the stream (2-approximation)

Robust Algorithms


- ❖ Start many instances of the streaming algorithm at the beginning
- ❖ Use an instance of the algorithm but “freeze” the output
- ❖ Each time the next instance has value $(1 + O(\epsilon))$ more than the “frozen” output, use the next instance and “freeze” its output



- ❖ Example: Number of ones in the stream (2-approximation)

Robust Algorithms


- ❖ Start many instances of the streaming algorithm at the beginning
- ❖ Use an instance of the algorithm but “freeze” the output
- ❖ Each time the next instance has value $(1 + O(\epsilon))$ more than the “frozen” output, use the next instance and “freeze” its output



- ❖ Example: Number of ones in the stream (2-approximation)

Robust Algorithms


- ❖ Start many instances of the streaming algorithm at the beginning
- ❖ Use an instance of the algorithm but “freeze” the output
- ❖ Each time the next instance has value $(1 + O(\epsilon))$ more than the “frozen” output, use the next instance and “freeze” its output



- ❖ Example: Number of ones in the stream (2-approximation)

Robust Algorithms


- ❖ Start many instances of the streaming algorithm at the beginning
- ❖ Use an instance of the algorithm but “freeze” the output
- ❖ Each time the next instance has value $(1 + O(\epsilon))$ more than the “frozen” output, use the next instance and “freeze” its output



- ❖ Example: Number of ones in the stream (2-approximation)

Robust Algorithms


- ❖ Start many instances of the streaming algorithm at the beginning
- ❖ Use an instance of the algorithm but “freeze” the output
- ❖ Each time the next instance has value $(1 + O(\epsilon))$ more than the “frozen” output, use the next instance and “freeze” its output



- ❖ Example: Number of ones in the stream (2-approximation)

Robust Algorithms


- ❖ Start many instances of the streaming algorithm at the beginning
- ❖ Use an instance of the algorithm but “freeze” the output
- ❖ Each time the next instance has value $(1 + O(\epsilon))$ more than the “frozen” output, use the next instance and “freeze” its output



- ❖ Example: Number of ones in the stream (2-approximation)

Robust Algorithms

- ❖ Start many instances of the streaming algorithm at the beginning
- ❖ Use an instance of the algorithm but “freeze” the output
- ❖ Each time the next instance has value $(1 + O(\epsilon))$ more than the “frozen” output, use the next instance and “freeze” its output




- ❖ Example: Number of ones in the stream (2-approximation)
- ❖ Number of ones stream is at least 4 and at most 8
- ❖ 4 is a good approximation

Summary

- ❖ Sketch switching for robust algorithms uses $\frac{1}{\epsilon^2}$ space each time F_p increases by $(1 + \epsilon)$ and function increases $\frac{1}{\epsilon}$ times
-
- ❖ Smooth histogram for sliding window algorithms uses $\frac{1}{\epsilon^2}$ space each time F_p increases by $(1 + \epsilon)$ and function increases $\frac{1}{\epsilon}$ times for $p \in (0,1)$
 - ❖ Smooth histogram for sliding window algorithms uses $\frac{1}{\epsilon^2}$ space each time F_p increases by $(1 + \epsilon^p)$ and function increases $\frac{1}{\epsilon^p}$ times for $p \in (1,2)$


Intuition

❖ Do we really need to pay $\frac{1}{\epsilon^2}$ space each time F_p increases by $(1 + \epsilon)$?




Intuition

❖ Do we really need to pay $\frac{1}{\epsilon^2}$ space each time F_p increases by $(1 + \epsilon)$?



Intuition


❖ Do we really need to pay $\frac{1}{\epsilon^2}$ space each time F_p increases by $(1 + \epsilon)$?



Suppose: $F_p(u+v) - F_p(u) = \epsilon F_p(u)$

Intuition


- ❖ Do we really need to pay $\frac{1}{\epsilon^2}$ space each time F_p increases by $(1 + \epsilon)$?
- ❖ Only need constant factor approximation to $\epsilon F_p(u)$
- ❖ Only need constant factor approximation to $F_p(u + v) - F_p(u)$



Suppose: $F_p(u + v) - F_p(u) = \epsilon F_p(u)$


Sketch Stitching

- ❖ Suppose we want $F_p(u + v)$
- ❖ $F_p(u + v) = (F_p(u + v) - F_p(u)) + F_p(u)$




Sketch Stitching

❖ Suppose we want $F_p(u + v)$ and $v = v_1 + v_2 + \cdots + v_b$



Sketch Stitching

- ❖ Suppose we want $F_p(u + v)$ and $v = v_1 + v_2 + \dots + v_b$
- ❖ $F_p(u + v) = (F_p(u + v) - F_p(u)) + F_p(u)$



Sketch Stitching


- ❖ Suppose we want $F_p(u + v)$ and $v = v_1 + v_2 + \cdots + v_b$
- ❖ $F_p(u + v) = (F_p(u + v) - F_p(u)) + F_p(u)$
- ❖ $F_p(u + v) = \left(F_p(u + v_1 + \cdots + v_b) - F_p(u + v_1 + \cdots + v_{b-1}) \right) + \left(F_p(u + v_1 + \cdots + v_{b-1}) - F_p(u + v_1 + \cdots + v_{b-2}) \right) + \cdots + \left(F_p(u + v_1) - F_p(u) \right) + F_p(u)$

Sketch Stitching

- ❖ Suppose we want $F_p(u + v)$ and $v = v_1 + v_2 + \cdots + v_b$
- ❖ $F_p(u + v) = (F_p(u + v) - F_p(u)) + F_p(u)$
- ❖ $F_p(u + v) = \left(F_p(u + v_1 + \cdots + v_b) - F_p(u + v_1 + \cdots + v_{b-1}) \right) + \left(F_p(u + v_1 + \cdots + v_{b-1}) - F_p(u + v_1 + \cdots + v_{b-2}) \right) + \cdots + \left(F_p(u + v_1) - F_p(u) \right) + F_p(u)$

Sketch Stitching

$$\begin{aligned} \blacklozenge F_p(u + v) &= \left(F_p(u + v_1 + \cdots + v_b) - F_p(u + v_1 + \cdots + v_{b-1}) \right) + \\ &\quad \left(F_p(u + v_1 + \cdots + v_{b-1}) - F_p(u + v_1 + \cdots + v_{b-2}) \right) + \cdots + \\ &\quad \left(F_p(u + v_1) - F_p(u) \right) + F_p(u) \end{aligned}$$




Granularity Change

- ❖ Set each difference to be exponentially decreasing
- ❖
$$F_p(u + v) = \left(F_p(u + v_1 + \cdots + v_b) - F_p(u + v_1 + \cdots + v_{b-1}) \right) + \left(F_p(u + v_1 + \cdots + v_{b-1}) - F_p(u + v_1 + \cdots + v_{b-2}) \right) + \cdots + \left(F_p(u + v_1) - F_p(u) \right) + F_p(u)$$
- ❖
$$F_p(u + v_1 + \cdots + v_b) - F_p(u + v_1 + \cdots + v_{b-1}) = \frac{1}{2^b} F_p(u)$$

Granularity Change

$$\blacklozenge F_p(u + v_1 + \cdots + v_b) - F_p(u + v_1 + \cdots + v_{b-1}) = \frac{1}{2^b} F_p(u)$$



Previously:





Granularity Change

- ❖ $F_p(u + v_1 + \cdots + v_b) - F_p(u + v_1 + \cdots + v_{b-1}) = \frac{1}{2^b} F_p(u)$
- ❖ Just need $2^b \epsilon$ -approximation to $F_p(u + v_1 + \cdots + v_b) - F_p(u + v_1 + \cdots + v_{b-1})$
- ❖ **Intuition:** pay $\frac{1}{\epsilon^2}$ for $(1 + \epsilon)$ -approximation
- ❖ Hope is to use space $\frac{1}{2^{2b} \epsilon^2}$ for $2^b \epsilon$ -approximation

Framework

- ❖ Algorithms simultaneously running for each granularity
- ❖ Want space $\frac{1}{2^{2b}\epsilon^2}$ for granularity $\frac{1}{2^b} F_p(u)$
- ❖ Need 2^b instances for granularity $\frac{1}{2^b} F_p(u)$



- ❖ Total space $\sum \frac{1}{2^{2b}\epsilon^2} = \frac{1}{\epsilon^2}$

Format

- ❖ Part 1: Background
- ❖ Part 2: Framework
- ❖ Part 3: Difference Estimators

Questions?



Difference Estimator

❖ If $F_p(u + v) - F_p(u) = 2^b \epsilon F_p(u)$, does there exist algorithm that approximates the difference with space $\frac{1}{2^{2b} \epsilon^2}$?



❖ **Definition:** $F_p(u + v) - F_p(u) = \gamma F_p(u)$, output an estimate to the difference with additive approximation $\epsilon F_p(u)$



Difference Estimator

- ❖ **Definition:** $F_p(u + v) - F_p(u) = \gamma F_p(u)$, output an estimate to the difference with additive approximation $\epsilon F_p(u)$
- ❖ F is generally non-linear
- ❖ **Ex:** $F_p(u + v) = \frac{1}{\epsilon^4}$, $F_p(u + v) - F_p(u) = 1$
- ❖ $(1 + \epsilon)$ approximations to $F_p(u + v)$ and $F_p(u)$ give multiplicative approximation to the difference but use space $\frac{1}{\epsilon^2}$
- ❖ Constant factor approximations to $F_p(u + v)$ and $F_p(u)$ do not give additive approximation $\epsilon F_p(u)$ to the difference



Our Results: Difference Estimators

- ❖ Space $\tilde{O}\left(\frac{\gamma}{\epsilon^2} \log n\right)$ algorithm for F_0
- ❖ Space $\tilde{O}\left(\frac{\gamma^{2/p}}{\epsilon^2} \log n\right)$ algorithm for F_p with $p \in (0, 2]$
- ❖ Space $\tilde{O}\left(\frac{\gamma}{\epsilon^2} n^{1-2/p}\right)$ algorithm for F_p with integer $p > 2$

F_2 Difference Estimator

- ❖ **Definition:** $F_2(u + v) - F_2(u) = \gamma F_2(u)$, output an estimate to the difference with additive approximation $\epsilon F_2(u)$
- ❖ $F_2(u + v) - F_2(u) = \langle u + v, u + v \rangle - \langle u, u \rangle = 2\langle u, v \rangle + \langle v, v \rangle^2$
- ❖ **Inner product property:** $(1 + \epsilon)$ -approximations to $\|u\|_2$ and $\|v\|_2$ gives an $\epsilon\|u\|_2\|v\|_2$ additive approximation to $\langle u, v \rangle$
- ❖ $\gamma F_2(u) \geq \langle v, v \rangle^2$ implies $2\langle u, v \rangle \leq 2\|u\|_2\|v\|_2 \leq 2\sqrt{\gamma} F_2(u)$
- ❖ Just need $\frac{\epsilon}{\sqrt{\gamma}}$ multiplicative approximation: $\tilde{O}\left(\frac{\gamma}{\epsilon^2} \log n\right)$ space!


F_2 Difference Estimator

- ❖ **Difference estimator:** Maintain $\left(1 + \frac{\epsilon}{\sqrt{\gamma}}\right)$ -approximations to $F_2(u + v)$ and $F_2(u)$ using AMS sketch
- ❖ For $F_2(u + v) - F_2(u) = \gamma F_2(u)$, difference of the outputs is an additive approximation $\epsilon F_2(u)$ to $F_2(u + v) - F_2(u)$
- ❖ Space $\tilde{O}\left(\frac{\gamma}{\epsilon^2} \log n\right)$ algorithm for difference estimator

Challenges for F_p Difference Estimators

- ❖ F_p difference estimator: Use p -stable random variables for $p \leq 2$?
- ❖ $\langle Z_p, f \rangle$ where Z_p has entries drawn from p -stable distribution

$$p(x) \sim \frac{1}{1 + |x|^{p+1}}$$

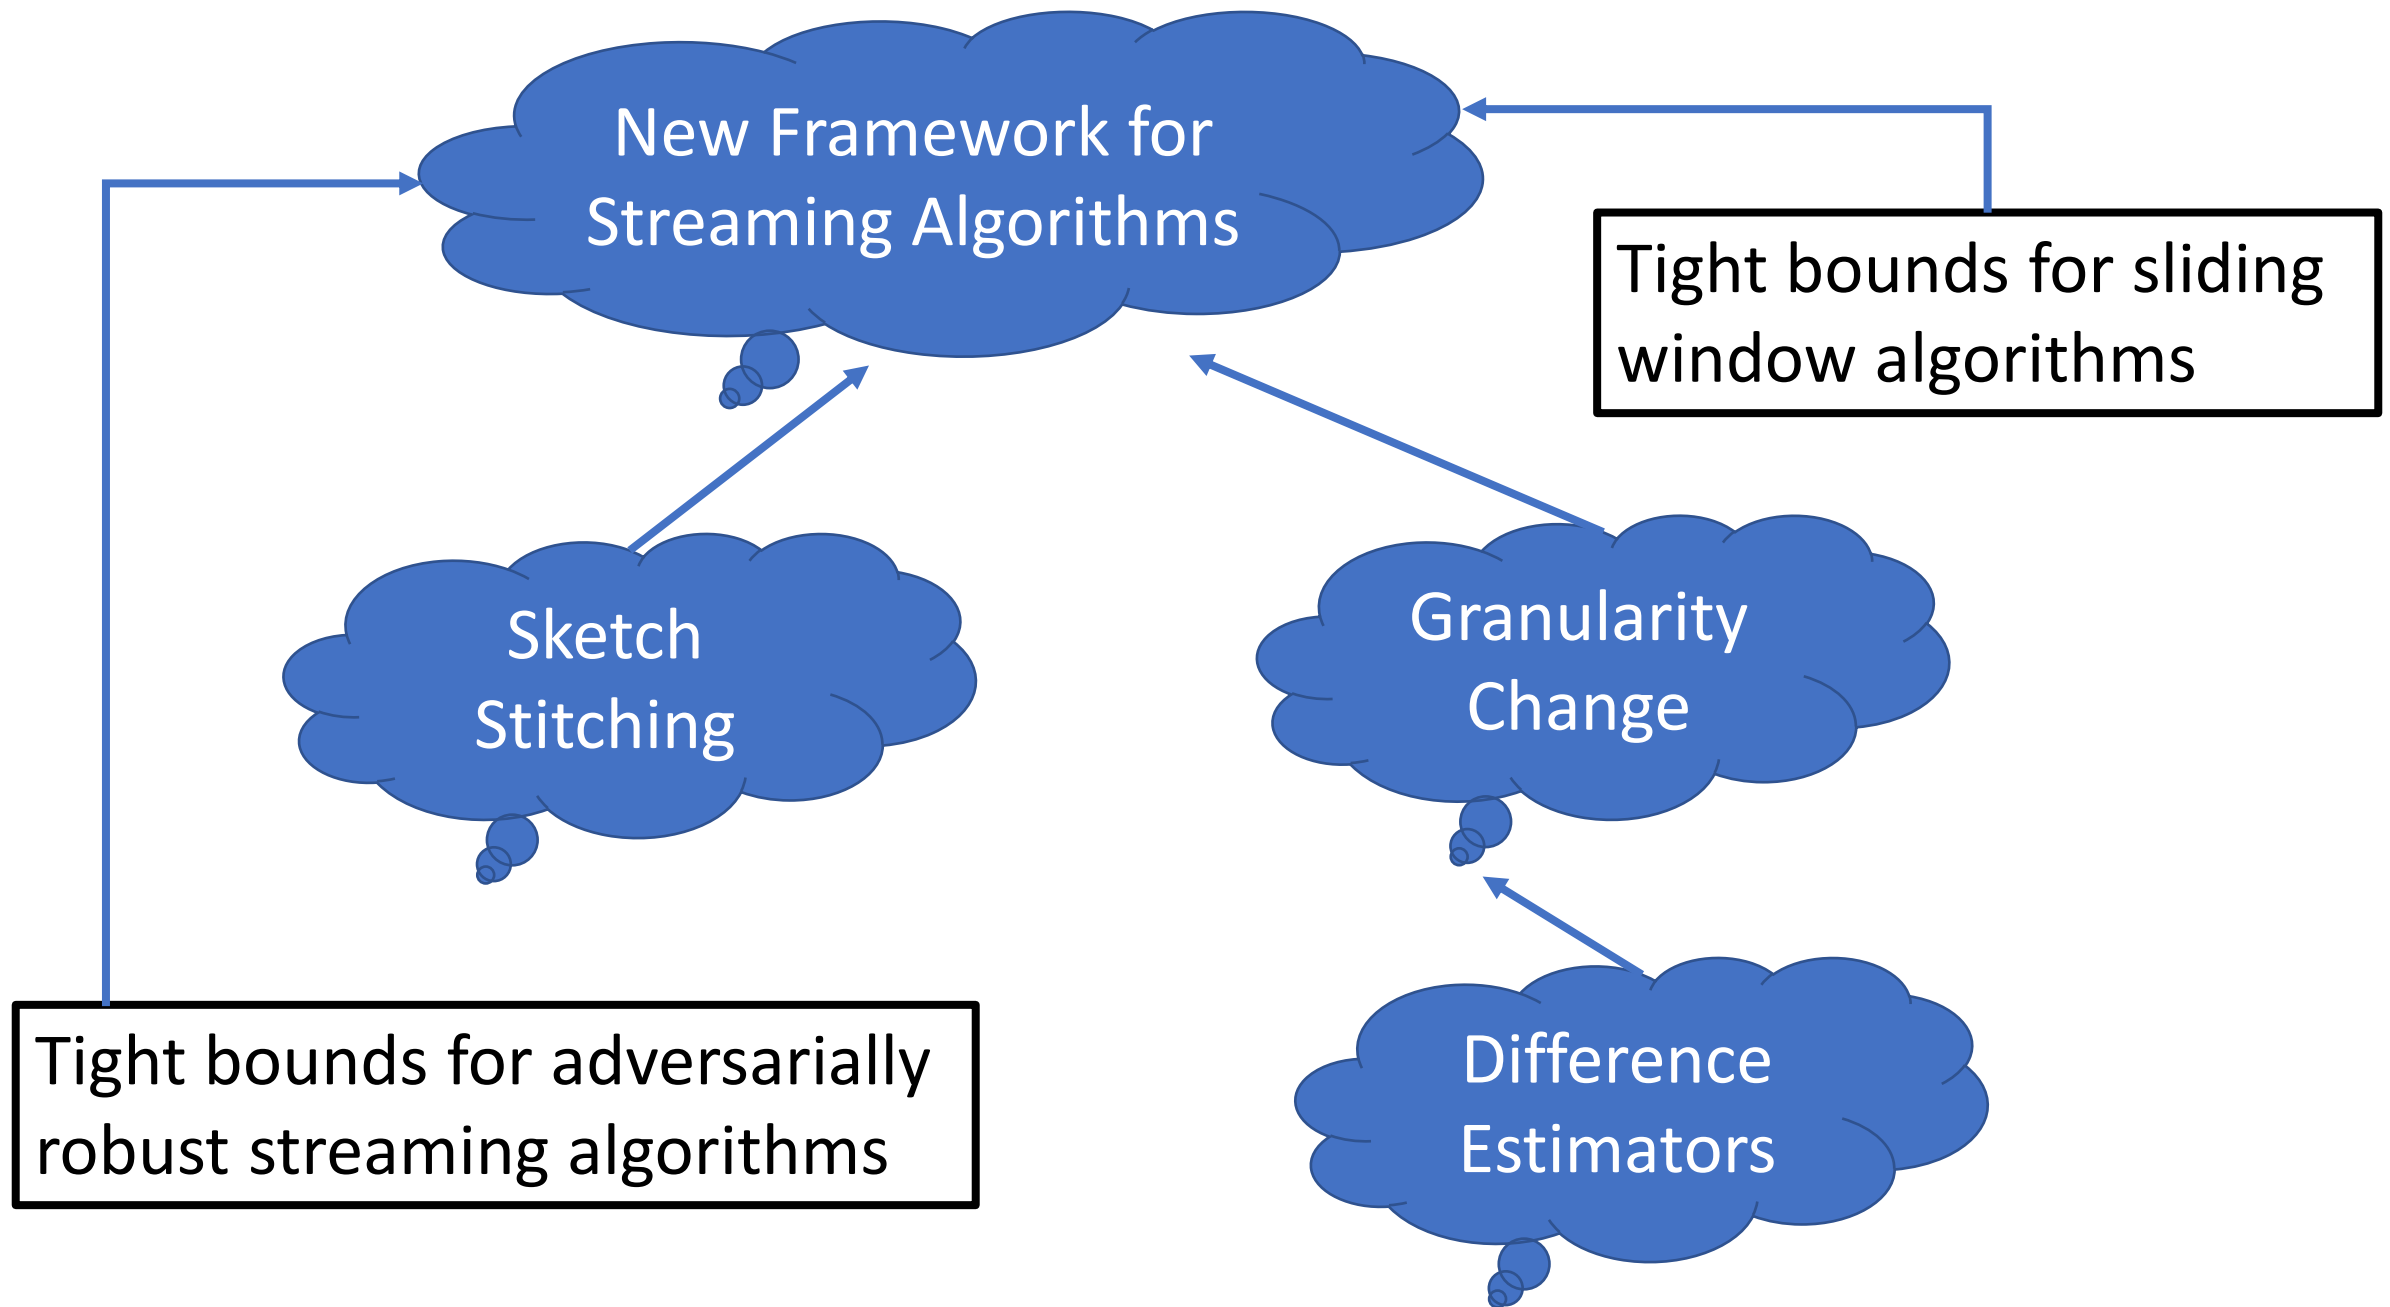


F_p Difference Estimators

- ❖ $Z = \text{median}\langle Z_p, f \rangle$
- ❖ How to analyze median of each estimate of $F_p(u + v) - F_p(u)$?
- ❖ Use Li's geometric mean algorithm [Li08]
- ❖ Take the geometric mean of 3 inner products $\langle Z_p, f \rangle$
- ❖ Take the average of $O\left(\frac{1}{\epsilon^2}\right)$ geometric means

F_p Difference Estimators

- ❖ **Difference estimator:** Maintain $\left(1 + \frac{\epsilon}{\gamma^{1/p}}\right)$ -approximations to $F_p(u + v)$ and $F_p(u)$ using Li's geometric mean estimator
- ❖ $A(u + v) - A(u) \sim \langle p_1, v \rangle^{p/3} \langle p_2, v \rangle^{p/3} \langle p_3, v \rangle^{p/3} + \langle p_1, u \rangle^{p/3} \langle p_2, v \rangle^{p/3} \langle p_1, v \rangle^{p/3} + \langle p_1, u \rangle^{p/3} \langle p_2, u \rangle^{p/3} \langle p_1, v \rangle^{p/3} + \dots$
- ❖ Each summand has $\langle p_1, v \rangle^{p/3}$ term, which has much smaller variance



Future Directions

- ❖ Additional applications of difference estimators, e.g, general $p > 2$?
- ❖ Uses of differential privacy for adaptive data analysis
- ❖ Algorithms robust to white-box adversaries?





Challenges for F_p Difference Estimators

- ❖ F_p difference estimator: Generalization of inner products for $p > 2$?
- ❖ Variance can be much larger!
- ❖ Use heavy-hitter algorithm to explicitly track “heavy” elements
- ❖ Use L_2 sampling algorithm with $n^{1-2/p}$ buckets to sample “light” elements

General Challenges

- ❖ Use known structural results from chaining to remove $\log n$ factor in difference estimator
- ❖ Avoids typical Chernoff + union bound argument by considering the expected supremum of a process, “strong tracking”
- ❖ Use suffix argument to remove $\log n$ factor in framework
- ❖ Adaptation to sliding window model

Robust vs. Sliding Window




Literature

- ❖ **Tight Bounds for Adversarially Robust Streams and Sliding Windows via Difference Estimators**, David P. Woodruff, Samson Zhou (FOCS 2021)
- ❖ **Adversarial Robustness of Streaming Algorithms through Importance Sampling**, Vladimir Braverman, Avinatan Hassidim, Yossi Matias, Mariano Schain, Sandeep Silwal, Samson Zhou (NeurIPS 2021)
- ❖ **The White-Box Adversarial Data Stream Model**, Miklós Ajtai, Vladimir Braverman, T.S. Jayram, Sandeep Silwal, Alec Sun, David P. Woodruff, Samson Zhou (PODS 2022)

Sliding Window Algorithms

- ❖ Suppose we are trying to approximate some given function
 1. Suppose we have a streaming algorithm for this function
 2. Suppose this function is “smooth”: If $f(B)$ is a “good” approximation to $f(A)$, then $f(B \cup C)$ will always be a “good” approximation to $f(A \cup C)$.




- ❖ Smooth histogram framework [\[BravermanOstrovsky07\]](#) gives a sliding window algorithm for this function

Smooth Histogram

- ❖ Suppose we are trying to approximate some given function
- ❖ Smooth histogram framework [BO07] gives a sliding window algorithm for this function
- ❖ Start a new instance of the streaming algorithm (along with existing instances) each time a new element arrives
- ❖ Each time there are three instances that report “close” values, delete the middle one
- ❖ Use different checkpoints to “sandwich” the sliding window

Smooth Histogram

- ❖ Start a new instance of the streaming algorithm (along with existing instances) each time a new element arrives
- ❖ Each time there are three instances that report “close” values, delete the middle one
- ❖ Use different checkpoints to “sandwich” the sliding window



- ❖ Example: Number of ones in sliding window (2-approximation)

Smooth Histogram


- ❖ Start a new instance of the streaming algorithm (along with existing instances) each time a new element arrives
- ❖ Each time there are three instances that report “close” values, delete the middle one
- ❖ Use different checkpoints to “sandwich” the sliding window



- ❖ Example: Number of ones in sliding window (2-approximation)

Smooth Histogram


- ❖ Start a new instance of the streaming algorithm (along with existing instances) each time a new element arrives
- ❖ Each time there are three instances that report “close” values, delete the middle one
- ❖ Use different checkpoints to “sandwich” the sliding window



- ❖ Example: Number of ones in sliding window (2-approximation)

Smooth Histogram


- ❖ Start a new instance of the streaming algorithm (along with existing instances) each time a new element arrives
- ❖ Each time there are three instances that report “close” values, delete the middle one
- ❖ Use different checkpoints to “sandwich” the sliding window



- ❖ Example: Number of ones in sliding window (2-approximation)

Smooth Histogram


- ❖ Start a new instance of the streaming algorithm (along with existing instances) each time a new element arrives
- ❖ Each time there are three instances that report “close” values, delete the middle one
- ❖ Use different checkpoints to “sandwich” the sliding window



- ❖ Example: Number of ones in sliding window (2-approximation)

Smooth Histogram


- ❖ Start a new instance of the streaming algorithm (along with existing instances) each time a new element arrives
- ❖ Each time there are three instances that report “close” values, delete the middle one
- ❖ Use different checkpoints to “sandwich” the sliding window



- ❖ Example: Number of ones in sliding window (2-approximation)

Smooth Histogram


- ❖ Start a new instance of the streaming algorithm (along with existing instances) each time a new element arrives
- ❖ Each time there are three instances that report “close” values, delete the middle one
- ❖ Use different checkpoints to “sandwich” the sliding window



- ❖ Example: Number of ones in sliding window (2-approximation)

Smooth Histogram


- ❖ Start a new instance of the streaming algorithm (along with existing instances) each time a new element arrives
- ❖ Each time there are three instances that report “close” values, delete the middle one
- ❖ Use different checkpoints to “sandwich” the sliding window



- ❖ Example: Number of ones in sliding window (2-approximation)

Smooth Histogram


- ❖ Start a new instance of the streaming algorithm (along with existing instances) each time a new element arrives
- ❖ Each time there are three instances that report “close” values, delete the middle one
- ❖ Use different checkpoints to “sandwich” the sliding window



- ❖ Example: Number of ones in sliding window (2-approximation)

Smooth Histogram

- ❖ Start a new instance of the streaming algorithm (along with existing instances) each time a new element arrives
- ❖ Each time there are three instances that report “close” values, delete the middle one
- ❖ Use different checkpoints to “sandwich” the sliding window



- ❖ Example: Number of ones in sliding window (2-approximation)
- ❖ Number of ones in sliding window is at least 4 and at most 7
- ❖ 4 is a good approximation

Granularity Change

- ❖ Set each difference to be exponentially decreasing
- ❖ $F_p(u + v) = \left(F_p(u + v_1 + \dots + v_b) - F_p(u + v_1 + \dots + v_{b-1}) \right) + \left(F_p(u + v_1 + \dots + v_{b-1}) - F_p(u + v_1 + \dots + v_{b-2}) \right) + \dots + \left(F_p(u + v_1) - F_p(u) \right) + F_p(u)$
- ❖ $F_p(u + v_1 + \dots + v_b) - F_p(u + v_1 + \dots + v_{b-1}) = \frac{1}{2^b} F_p(u)$
- ❖ Just need $2^b \epsilon$ -approximation to $F_p(u + v_1 + \dots + v_b) - F_p(u + v_1 + \dots + v_{b-1})$
- ❖ Hope is to use space $\frac{1}{2^{2b} \epsilon^2}$