
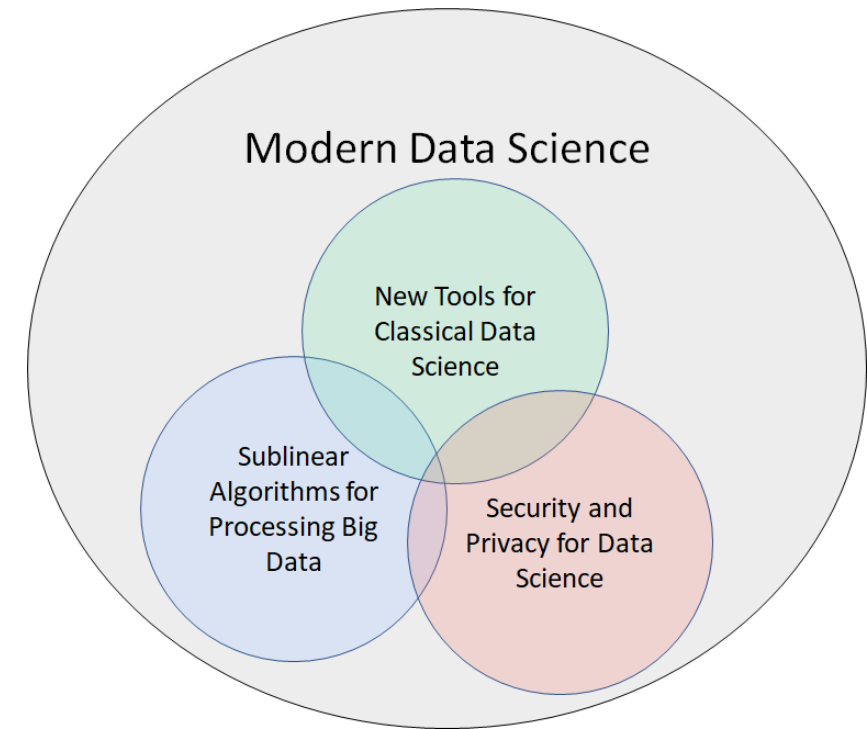



Theoretical Foundations of Modern Data Science

Samson Zhou



Background

- Post-doc at UC Berkeley / Rice
- **Previous**: Post-doc at Carnegie Mellon
- PhD in Computer Science from Purdue
- Dual Bachelors in Math, Computer Science from MIT

Berkeley
UNIVERSITY OF CALIFORNIA



Carnegie
Mellon
University





- **Research areas**: Data science, sublinear algorithms, security and privacy

Data Science











0







1 1 0 1 0 1
0 1 1 1 0 0



0
1









1	1	0	1	0	1
0	1	1	1	0	0
0	1	0	1	1	1
0	1	1	0	1	0
1	0	1	1	1	1
0	1	1	1	0	1
0	0	1	0	1	1



0
1
1
0
0
1
0







new ideas
science
engineering



Mark Rober ✓

@MarkRober

23.3M subscribers

< SHORTS LIVE PLAYLISTS COMMUNITY STORE CHANNELS ABOUT





Stealing Baseball Signs with a Phone (Machine Learning)

Mark Rober ✓ 24M views • 3 years ago

I always sucked at baseball... until now... ok, I still probably suck. Go subscribe to Jabril's channel!!!
<https://www.youtube.com/channel/UCQALLLeQPozdZC4JNUboVEUg> Simple app (it's actually...

CC





Search

About 13,400,000 results (0.44 seconds)

Warning: This page is a tool for AdWords advertisers to test their ads. For full Google functionality [return to the Google homepage](#).

Web

Images

Maps

Videos

News

More

Auckland

Change location

The web

Pages from New Zealand

More search tools

Ads related to chocolate gift baskets ⓘ

[Gift Baskets | BeautifulBaskets.co.nz](#)

www.beautifulbaskets.co.nz/

NZ's Favourite **Gift Baskets** Fast Nationwide NZ Delivery

[Chocolates Delivered Fast - Delicious Handmade NZ Chocolates.](#)

www.devonportchocolates.co.nz/

Browse Our Website & Order Online.

Devonport, 17 Wynyard Street, Auckland - 0800 002 462 - [Directions](#)

Chocolate Bars - Weddings & Special Occasions - Birthday Gifts - Corporate Gifts

[Great Gift Baskets Online - Stunning Hampers for all Occasions.](#)

www.wineplus.co.nz/gift-baskets

Fast Delivery NZ-wide - Buy Online.

[Chocolate Gift Baskets - Auckland Flowers and Gifts](#)

www.nzflower.co.nz/.../gift_baskets_chocolate_new_zealand_auckla...

Chocolate Gift Baskets: Specializing in Chocolate Baskets, Chocolate Gifts, Gourmet **Chocolate Gift Baskets**, Chocolate Lovers Gift Baskets. Easy and Secure ...

[Chocolate Gift Baskets | Bliss Baskets and Gifts](#)

www.blissgiftbaskets.co.nz/style/chocolate-gift-baskets

With their wide range of variety and versatility, **chocolate gift baskets** are a popular gift that is always appreciated by the recipient.

[Gift Ideas, Chocolate Bouquets | Edible Blooms NZ](#)

www.edibleblooms.co.nz/

Edible Blooms New Zealand offers a unique twist on flowers and **gift hampers**. Our range of **chocolate** bouquets, fresh fruit bouquets and gourmet **gift baskets** ...

[Gift Baskets for Chocolate Lovers::My Goodness Gift Baskets New ...](#)

Ads ⓘ

[Order Gift Baskets Online](#)

www.hamperbiz.co.nz/Gifts

Stunning Range of **Gifts**, Hampers
Wine, Flowers, Free Delivery.

[Gift Baskets Gourmet Food](#)

www.giftbarn.co.nz/

Fine Wine & **Chocolates**

Next day delivery New Zealand Wide

[Chocolate delivery](#)

www.edibleblooms.co.nz/

Chocolate Bouquets The Perfect **Gift**
Order New Zealand Wide Online Now

[Top Christmas Hampers](#)

www.champershampers.co.nz/

The highest quality food and wine
corporate hampers, Christmas hams

[NZ's Top Gift Baskets](#)

www.mygoodness.co.nz/

Quality, Stylish Hampers from Award
Winning, Great Service Kiwi Company


[Gift Baskets Delivered NZ](#)

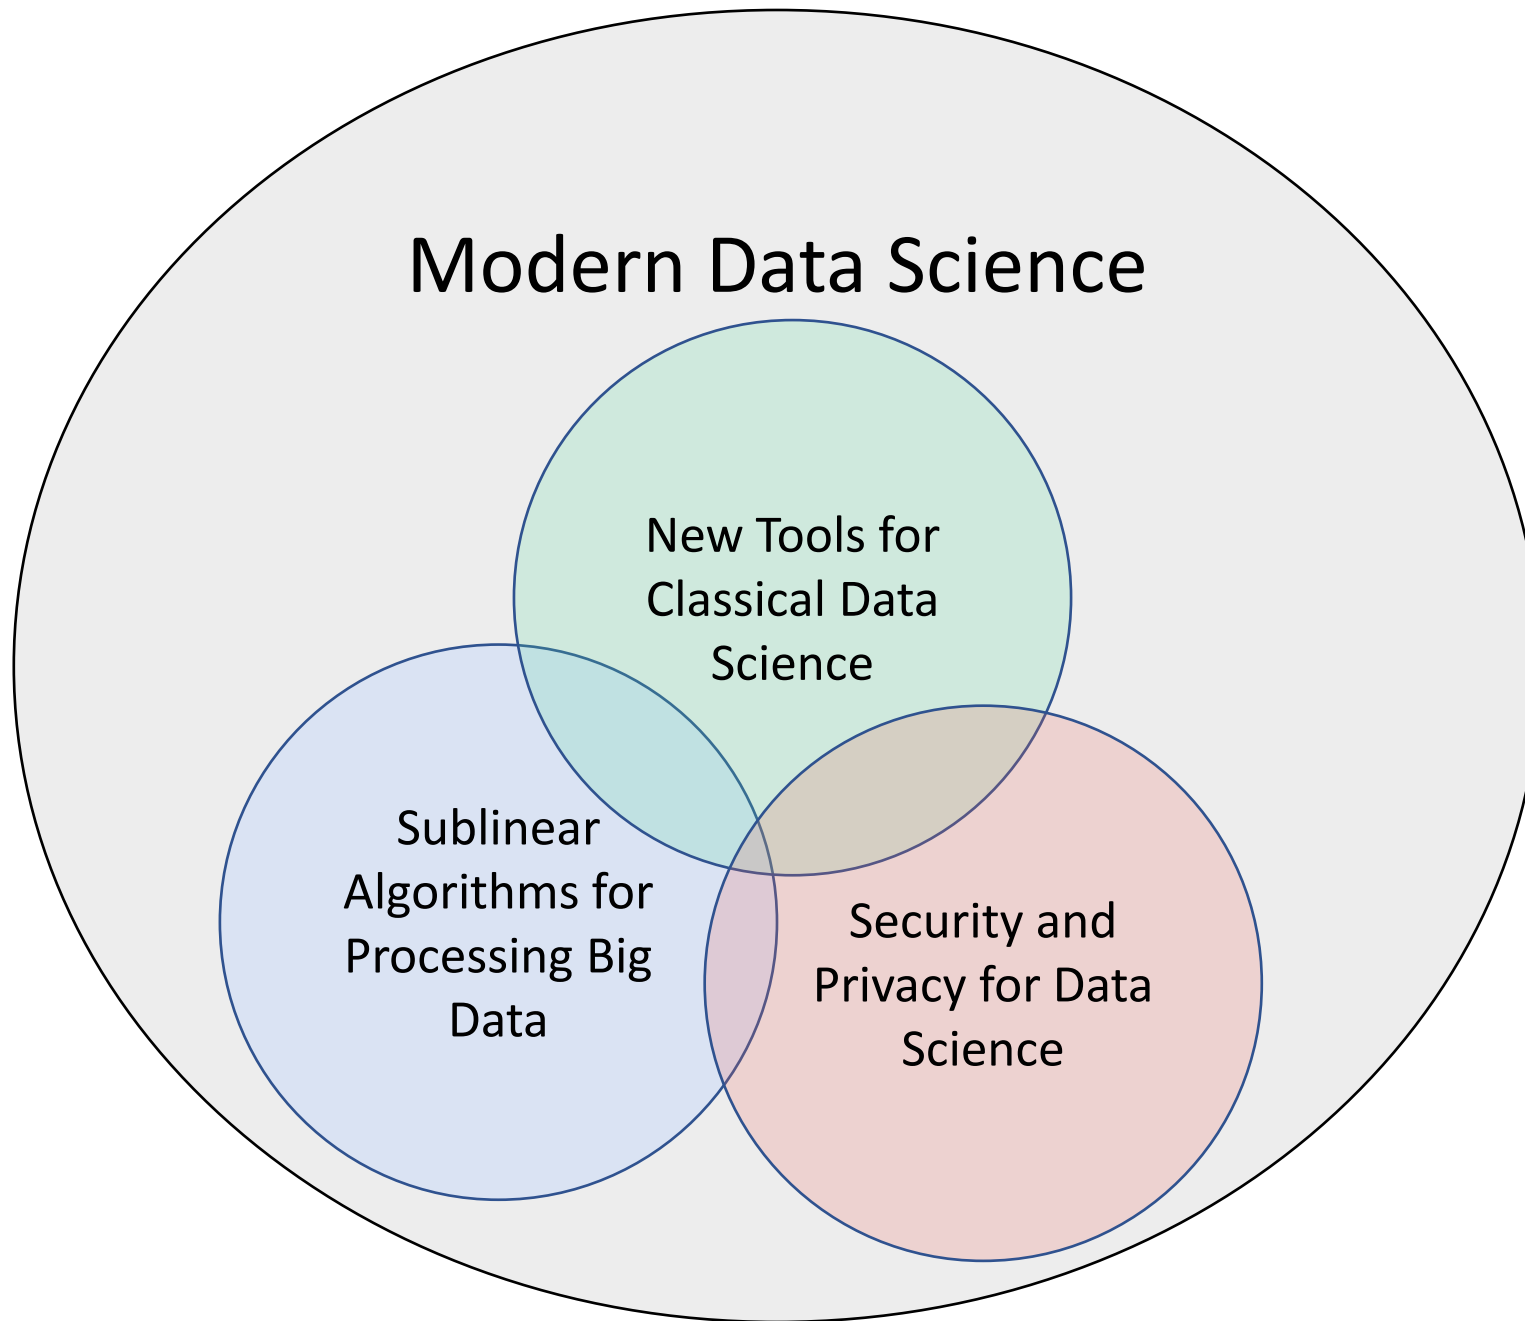
www.raptaboutgifts.co.nz/Gifts

Huge Selection of **Gift Baskets**
Browse our Online Store Now!

[Fine Chocolate Delivery](#)


Google Ad Revenue (2001-2021, billion USD)





Evolving Demands

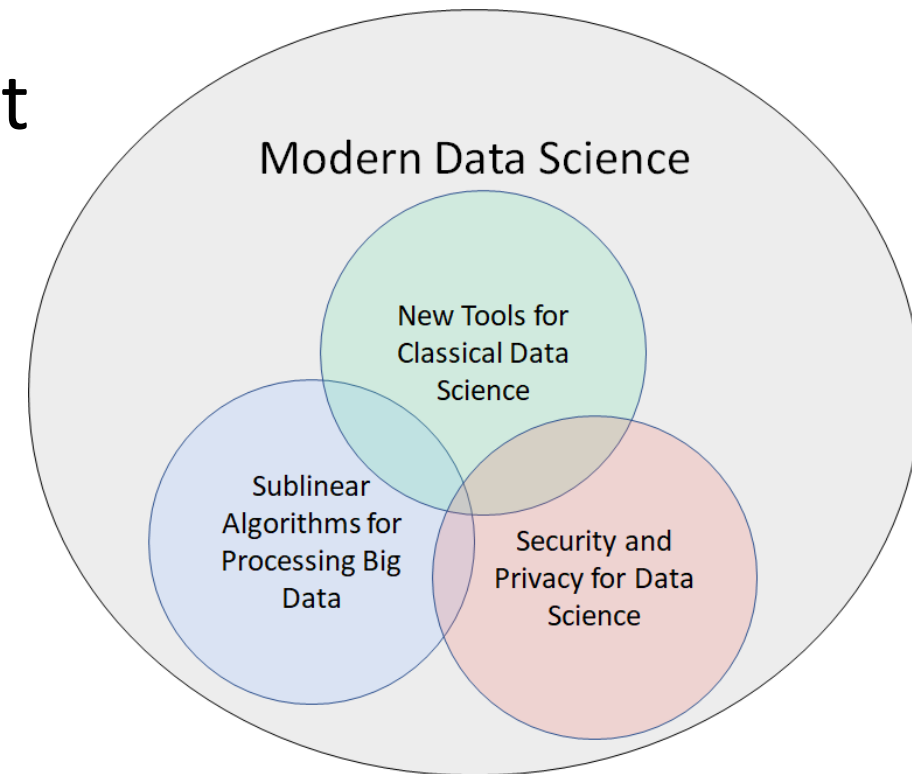
- Sublinear-time or sublinear-space algorithms
- Incorporation of advice
- Security and privacy
- Robustness to noise or adversarial input
- Ability to handle time-sensitive data



New Tools for Classical Data Science

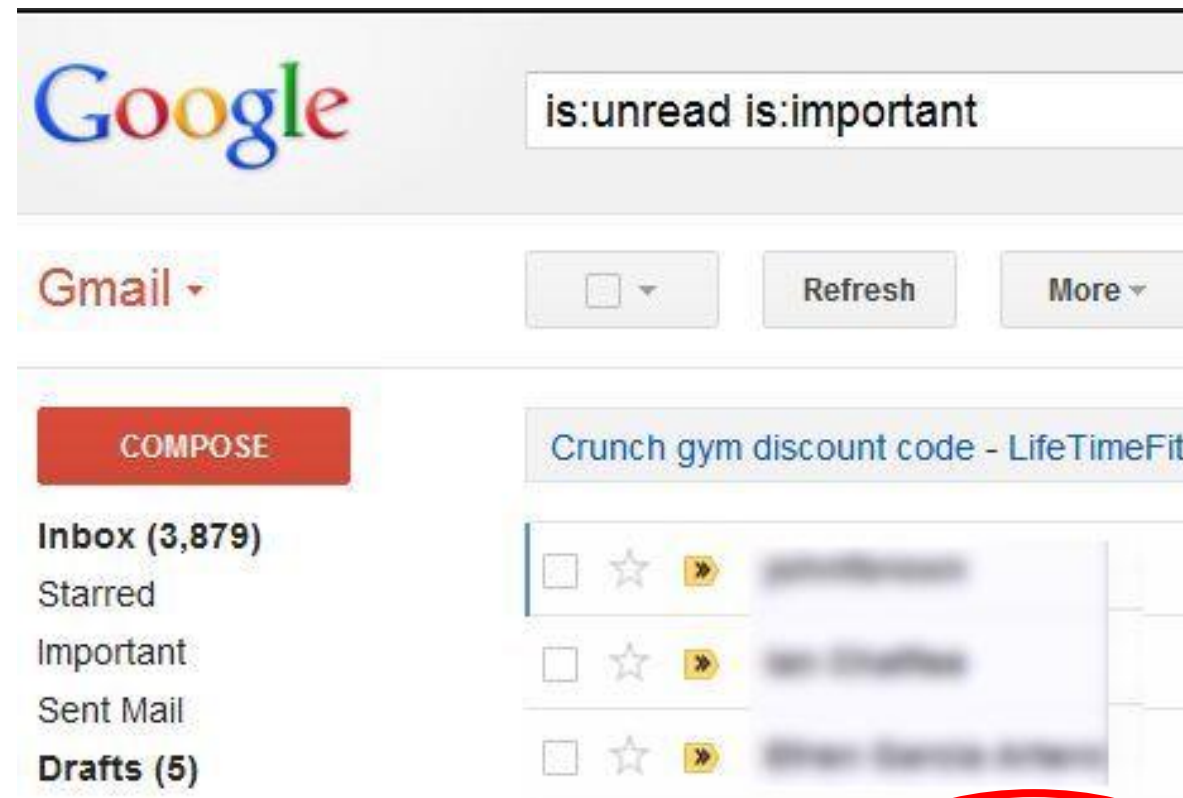
Evolving Demands

- Sublinear-time or sublinear-space algorithms
- Incorporation of advice
- Security and privacy
- Robustness to noise or adversarial input
- Ability to handle time-sensitive data

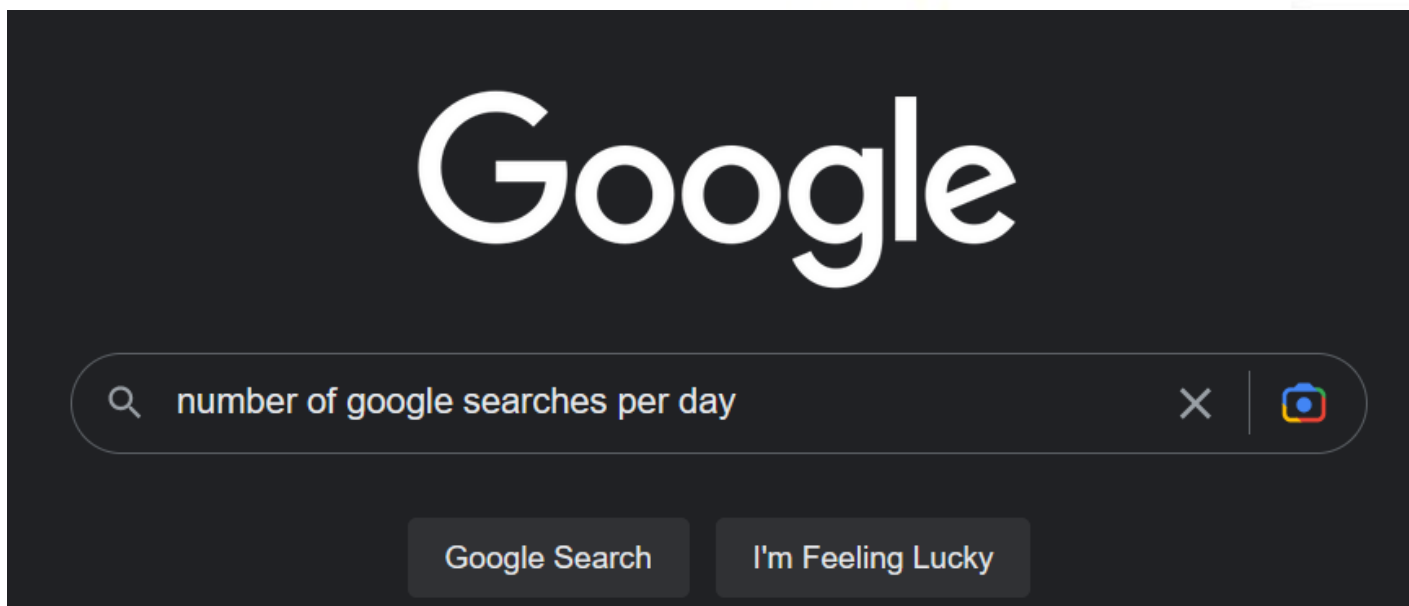





3 billion
monthly
active users



330 billion
daily e-mails

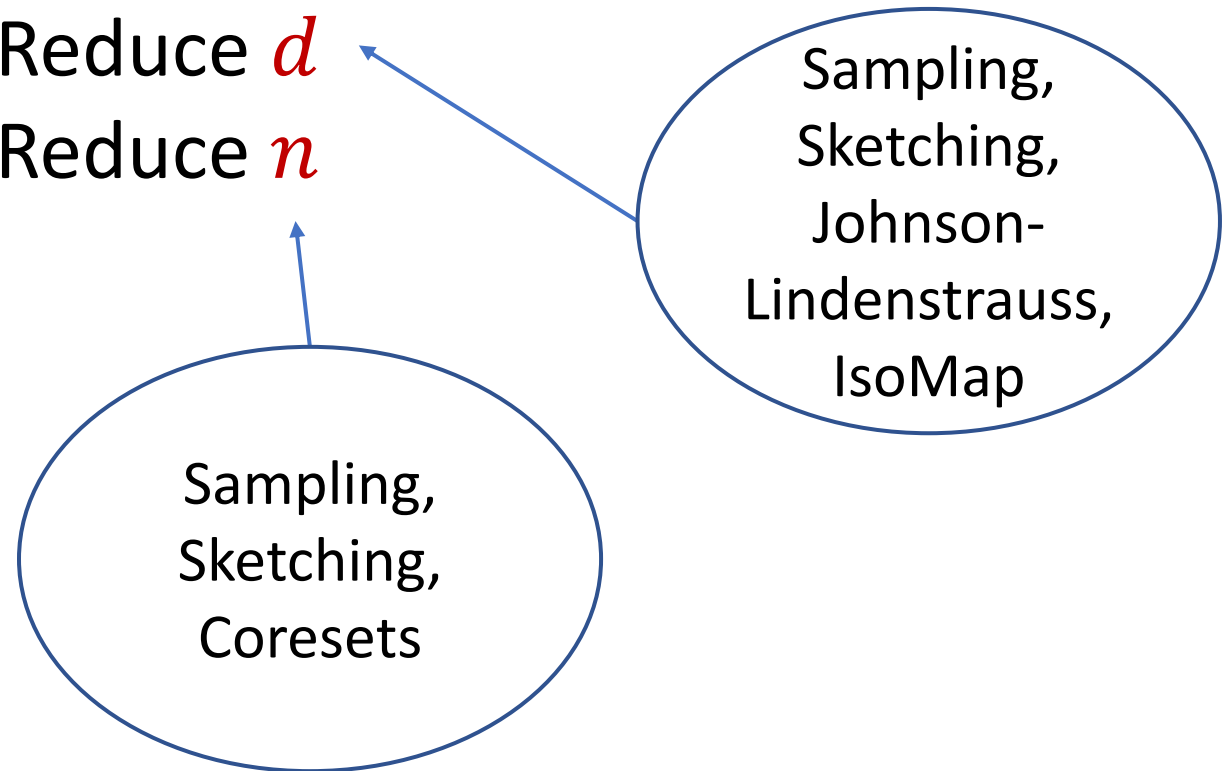


8.5 billion
daily Google
searches



Dimensionality Reduction

- Reduce d
- Reduce n




- Application depends on task!

New Dimensionality Reduction Results

- Dimensionality reduction for Wasserstein barycenter [ISZ21]
- Sketching for valuation functions, e.g., additive, submodular, Lipschitz [YZ19]
- Kernel density estimation on kernel matrices [BKISZ22]

Coreset

- Subset A' of representative rows of A for a given task with “score” function f
- $f(A, \cdot) \approx f(A', \cdot)$




Coreset

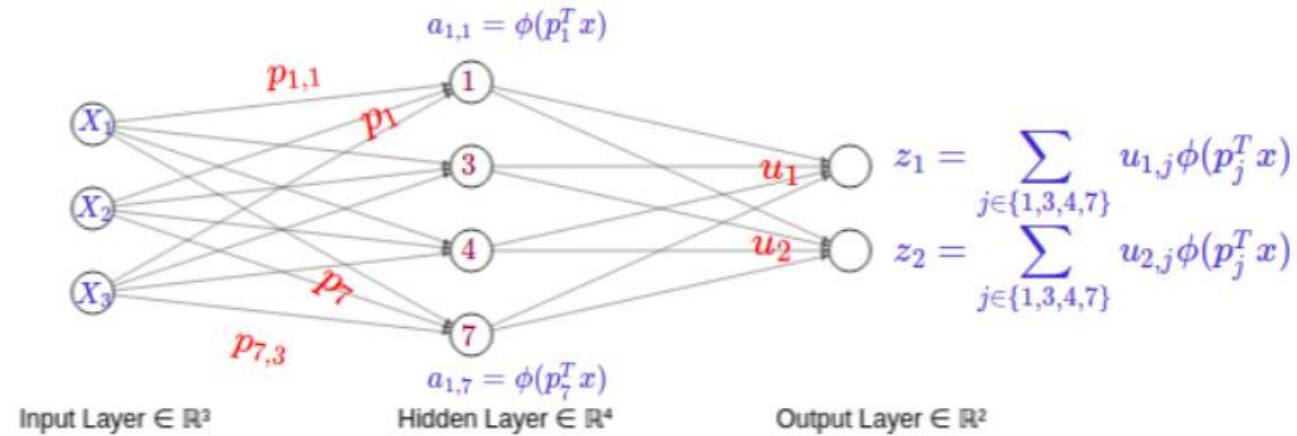
- Lots of different constructions with various tradeoffs
 - Size n' vs. accuracy
 - Size n' vs. computation time
 - Size n' vs. interpretability
 - Average-case vs. worst-case performance

New Coreset Constructions

- Coresets for data-independent neural pruning [MOBZF20]
- Coresets for projective clustering [TWZBF22]
- Coresets for regression on structured matrices [MMMWZF22]
- Online coresets for linear algebra [BDMMUWZ20]
- General coreset construction framework through sensitivity sampling [BFLSZ21]



(a)



(b)

[MOBZF20]

Network	Error(%)	# Parameters	Compression Ratio
LeNet-300-100	2.16	267K	90%
LeNet-300-100 Pruned	2.03	26K	
VGG-16	8.95	1.4M	75%
VGG-16 Pruned	8.16	350K	

Table 2: Empirical evaluations of our coresets on existing architectures for MNIST and CIFAR-10. Note the improvement of accuracy in both cases!

[MOBZF20]

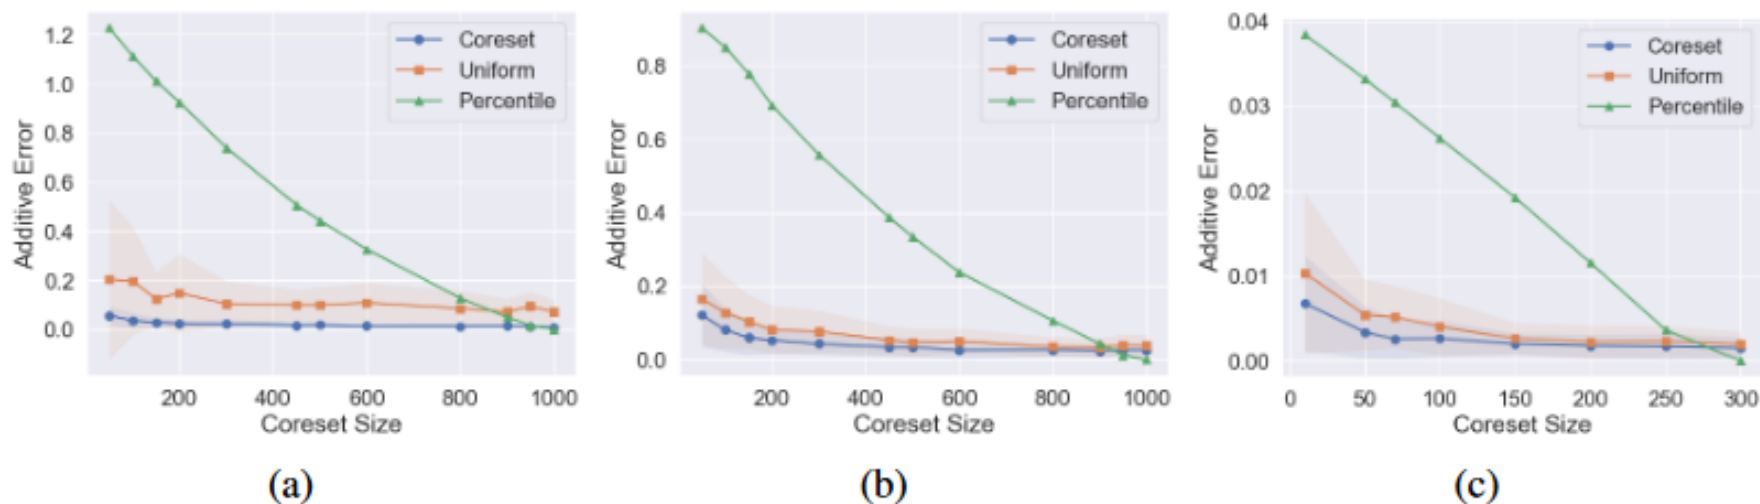




Figure 2: Approximation error of a single neuron on MNIST dataset across different coreset sizes. The weights of the points in (a) are drawn from the Gaussian distribution, in (b) from the Uniform distribution and in (c) we used the trained weights from LeNet-300-100.

Evolving Demands

- Sublinear-time or sublinear-space algorithms
- **Incorporation of advice**
- Security and privacy
- Robustness to noise or adversarial input
- Ability to handle time-sensitive data








29	Data-Driven Algorithm Design	626
	<i>Maria-Florina Balcan</i>	
29.1	Motivation and Context	626
29.2	Data-Driven Algorithm Design via Statistical Learning	628
29.3	Data-Driven Algorithm Design via Online Learning	639
29.4	Summary and Discussion	644
30	Algorithms with Predictions	646
	<i>Michael Mitzenmacher and Sergei Vassilvitskii</i>	
30.1	Introduction	646
30.2	Counting Sketches	649
30.3	Learned Bloom Filters	650
30.4	Caching with Predictions	652
30.5	Scheduling with Predictions	655
30.6	Notes	660

ABOUT



scheduling

scheduling

differential privacy

online

online


routing

SDP

Learning-Augmented Clustering

- **Goal:** Given dataset P in d dimensions, output a set C of k centers to minimize

$$\sum_{p \in P} \min_{c \in C} \|p - c\|_2^2$$



Learning-Augmented Clustering


- **Goal:** Given dataset P in d dimensions, output a set C of k centers to minimize

$$\sum_{p \in P} \min_{c \in C} \|p - c\|_2^2$$

- NP-hard to even approximate within a factor of 1.07 [CC20, LSW17]
- **Beyond worst case:** Clustering on similar inputs or inputs with auxiliary information
- ML can guide the clustering!
- **Our main message:** We can avoid worst case with advice!

Predictor

- Suppose Π outputs noisy labels according to a $(1 + \alpha)$ approximate clustering C and error rate $\lambda \leq \alpha$



What is the
label of x_1 ?

What is the
label of x_2 ?

x_1 belongs
to cluster 3

x_2 belongs
to cluster 7




Theoretical Guarantee

- Suppose Π outputs noisy labels according to a $(1 + \alpha)$ approximate clustering C and error rate $\lambda \leq \alpha$
- **Main result [EFSWZ22]**: Algorithm that outputs a $(1 + O(\alpha))$ approximate k -means clustering in nearly linear time
- “Predictions can overcome complexity hardness barriers!”

Algorithmic Intuition

- Not enough to blindly follow predictions!




- Optimal cost ≈ 0
- Predictor with arbitrary small error has large cost!


Algorithmic Intuition

- Not enough to blindly follow predictions!
- **Our approach:** Use ideas from robust mean estimation

$$(1 - \varepsilon)P$$



$$\varepsilon Q$$



Algorithmic Intuition

- Not enough to blindly follow predictions!
- **Our approach:** Use robust mean estimation



Algorithmic Intuition

- Not enough to blindly follow predictions!
- **Our approach:** Use robust mean estimation



Algorithmic Intuition

- Not enough to blindly follow predictions!
- **Our approach:** Use robust mean estimation
- We exploit the fact that cluster centers are means of points!


Algorithmic Intuition

- Not enough to blindly follow predictions!
- **Our approach:** Use robust mean estimation
- We exploit the fact that cluster centers are means of points!
- Introduces a new set of tools for k -means clustering

Experimental Results

- **Case Study:** Spectral clustering on graphs varying over time
- **Dataset:** Internet router graph varying over the course of a year
- **Methodology:** Compare to standard benchmarks while using various natural predictors, i.e., noisily perturb true labels and compare to baselines as function of error

Dataset: Internet router graph varying over the course of a year, $k = 10$



Conclusion: Our algorithm (using predictor) outperforms benchmarks such as k -means ++ for low error while staying competitive with high corruptions

My Recent Work on Data Science

- Learning-Augmented k-means Clustering (ICLR 2022)
- Fast Regression for Structured Inputs (ICLR 2022)
- New Coresets for Projective Clustering and Applications (AISTATS 2022)
- Dimensionality Reduction for Wasserstein Barycenter (NeurIPS 2021)
- Learning a Latent Simplex in Input Sparsity Time (ICLR 2021)
- Near Optimal Linear Algebra in the Online and Sliding Window Models (FOCS 2020)
- Data-Independent Neural Pruning via Coresets (ICLR 2020)
- Adversarially Robust Submodular Maximization under Knapsack Constraints (KDD 2019)

Questions?

Security and Privacy for Big Data

2017 Equifax Data Breach



“Equifax agreed to a \$700 million settlement over the privacy breach, but \$425 million of that was set aside to repay consumers as a restitution fund.”

YAHOO!



eHarmony®



LastPass*****

ebay




LinkedIn™




Evolving Demands

- Sublinear-time or sublinear-space algorithms
- Incorporation of advice
- **Security and privacy**
- Robustness to noise or adversarial input
- Ability to handle time-sensitive data




Graph Theory for Memory Hard Functions



Graph Theory for Memory Hard Functions

- Proved the security benefits of memory-hard functions using an economic marginal cost-revenue analysis [BHZ18]
- Showed the hardness of computing [BZ18] and approximating the computational cost [BLZ20] of graph pebbling
- Applied graph theory to cryptanalyze the computational complexity of memory-hard functions [BZ18, BRZ18]



census.gov:



Privacy & Confidentiality

Federal Law Protects Your Information. The U.S. Census Bureau is bound by [Title 13](#) of the United States Code. This law not only provides authority for the work we do, but also provides strong protection for the information we collect from individuals and businesses. As a result, the Census Bureau has one of the strongest confidentiality guarantees in the federal government.

It is against the law for any Census Bureau employee to disclose or publish any census or survey information that identifies an individual or business. This is true even for inter-agency communication: the FBI and other government entities do not have the legal right to access this information. In fact, when these protections have been challenged, Title 13's confidentiality guarantee has been upheld.

For more information about how the Census Bureau safeguards the data it collects, visit the agency's [Data Protection](#) and [Disclosure Avoidance Working Papers](#) Web sites.

Anonymizing Data

Age	Zip Code	Employer	Has Pet
56	77005	Apple	Yes
32	77005	Microsoft	No
71	77005	Amazon	Yes
44	77005	Petsmart	Yes
25	77005	Netflix	No
61	77005	Google	No

Anonymizing Data

Age	Zip Code	Employer	Has Pet
56	77005	Apple	Yes
32	77005	Microsoft	No
71	77005	Amazon	Yes
44	77005	Petsmart	Yes
25	77005	Netflix	No
61	77005	Google	No

Name	Age	Gender	Employer
Alice	56	Female	Apple
Bob	32	Male	Microsoft
Carol	71	Female	Amazon
Dale	44	Male	Petsmart
Erin	25	Female	Netflix
Fred	61	Male	Google

Reconstruction Attack

Name	Age	Zip Code	Gender	Employer	Has Pet
Alice	56	77005	Female	Apple	Yes
Bob	32	77005	Male	Microsoft	No
Carol	71	77005	Female	Amazon	Yes
Dale	44	77005	Male	Petsmart	Yes
Erin	25	77005	Female	Netflix	No
Fred	61	77005	Male	Google	No

Implications of the simulated attack

The Census Bureau believed in 2010 that it was necessary to coarsen geographic identifiers in microdata such that the minimum population in any published geography was at least 100,000 persons (Public-Use Microdata Areas).

Our simulated reconstruction-abetted re-identification attack demonstrated that the tabular summaries from the 2010 Census can be converted into a 100% microdata file with geographic precision to the census block-level.


Our simulated attack demonstrated that, depending on the quality of the external data used, between 52 and 179 million respondents to the 2010 Census can be correctly re-identified from the reconstructed microdata.

Stronger privacy protections, such as those in the 2020 Census Disclosure Avoidance System, are necessary to protect against reconstruction-abetted attacks.

Differential Privacy

- [DMNS06] Given $\varepsilon > 0$ and $\delta \in (0,1)$, a randomized algorithm $A: U^* \rightarrow Y$ is (ε, δ) -differentially private if, for every neighboring frequency vectors f and f' and for all $E \subseteq Y$,

$$\Pr[A(f) \in E] \leq e^\varepsilon \Pr[A(f') \in E] + \delta$$



Differential Privacy Result Highlights

- Framework for converting FPTAS/FPRAS algorithms to private FPTAS/FPRAS algorithms [BGMZ22]
- Privacy has a price in memory cost [DSWZ23]
- No overhead in communication cost for summation in the differentially oblivious shuffle model [GKNMZ22]

My Recent Work on Security and Privacy

- Private Data Stream Analysis for Universal Symmetric Norm Estimation (FORC 2022)
- On the Security of Proofs of Sequential Work in a Post-Quantum World (ITC 2021)
- Approximating Cumulative Pebbling Cost is Unique Games Hard (ITCS 2020)
- Computationally Data-Independent Memory Hard Functions (ITCS 2020)
- Data-Independent Memory Hard Functions: New Attacks and Stronger Constructions (CRYPTO 2019)
- Bandwidth-Hard Functions: Reductions and Lower Bounds (CCS 2018)
- On the Economics of Offline Password Cracking (Security and Privacy 2018)

Questions?

Sublinear Algorithms for Processing Big Data

Model #1: Streaming Model

- **Input:** Elements of an underlying data set S , which arrives sequentially
- **Output:** Evaluation (or approximation) of a given function
- **Goal:** Use space *sublinear* in the size m of the input S

1 0 1 1 1 0 0 1

Frequency Vector

- Given a set S of m elements from $[n]$, let f_i be the frequency of element i . (How often it appears)

$$1\ 1\ 2\ 1\ 2\ 1\ 1\ 2\ 3 \rightarrow [5, 3, 1, 0] := f$$

Heavy-Hitters (Frequent Items)

- Given a set S of m elements from $[n]$, let f_i be the frequency of element i . (How often it appears)
- Let L_2 be the norm of the frequency vector:

$$L_2 = \sqrt{f_1^2 + f_2^2 + \cdots + f_n^2}$$

- **Goal:** Given a set S of m elements from $[n]$ and a threshold ε , output the elements i such that $f_i > \varepsilon L_2$...and no elements j such that $f_j < \frac{\varepsilon}{16} L_2$
- **Motivation:** DDoS prevention, iceberg queries

Frequency Moments (L_p Norm)

- Given a set S of m elements from $[n]$, let f_i be the frequency of element i . (How often it appears)
- Let F_p be the frequency moment of the vector:

$$F_p = f_1^p + f_2^p + \cdots + f_n^p$$

- **Goal:** Given a set S of m elements from $[n]$ and an accuracy parameter ε , output a $(1 + \varepsilon)$ -approximation to F_p
- **Motivation:** Entropy estimation, linear regression

Distinct Elements (F_0 Estimation)

- Given a set S of m elements from $[n]$, let f_i be the frequency of element i . (How often it appears)
- Let F_0 be the frequency moment of the vector:

$$F_0 = |\{i : f_i \neq 0\}|$$

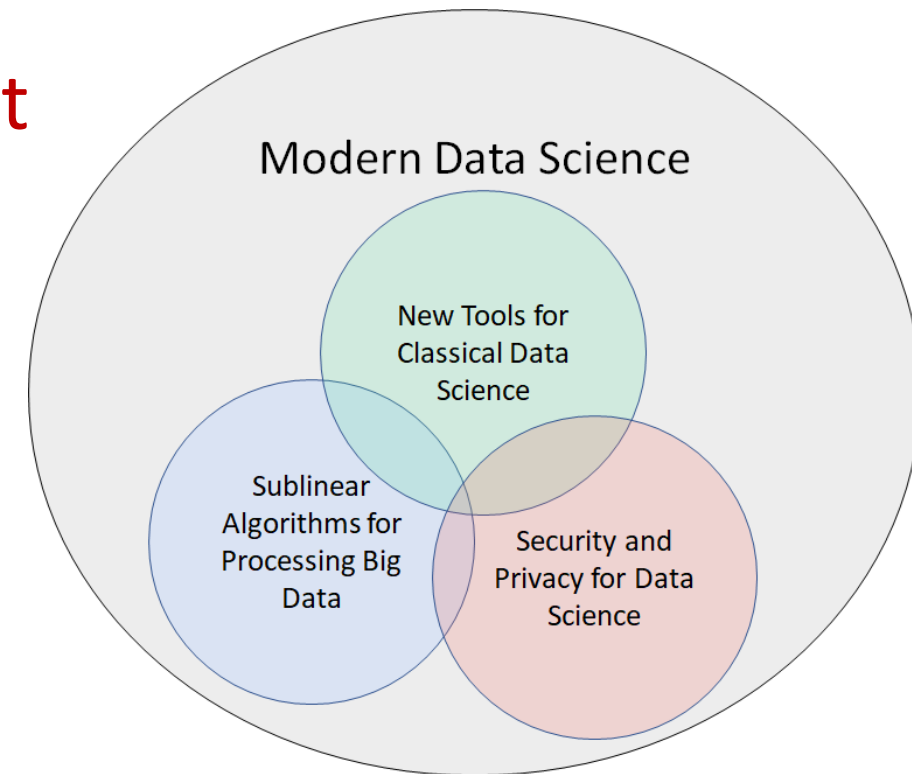
- **Goal:** Given a set S of m elements from $[n]$ and an accuracy parameter ε , output a $(1 + \varepsilon)$ -approximation to F_0
- **Motivation:** Traffic monitoring

$(1 + \varepsilon)$ -Approximation Streaming Algorithms

- $\Theta\left(\frac{1}{\varepsilon^2}\right)$ space-accuracy tradeoff for distinct elements, i.e., F_0 [KNW10, Blasiok20]
- $\Theta\left(\frac{1}{\varepsilon^2}\right)$ space-accuracy tradeoff for F_p with $p \in (0, 2]$ [BDN17]
- $\Theta\left(\frac{1}{\varepsilon^2}\right)$ space-accuracy tradeoff for F_p with $p > 2$ [Ganguly11, GW18]
- $\Theta\left(\frac{1}{\varepsilon^2}\right)$ space-accuracy tradeoff for L_2 -heavy hitters [BCINWW17]

Evolving Demands

- Sublinear-time or sublinear-space algorithms
- Incorporation of advice
- Security and privacy
- **Robustness to noise or adversarial input**
- Ability to handle time-sensitive data

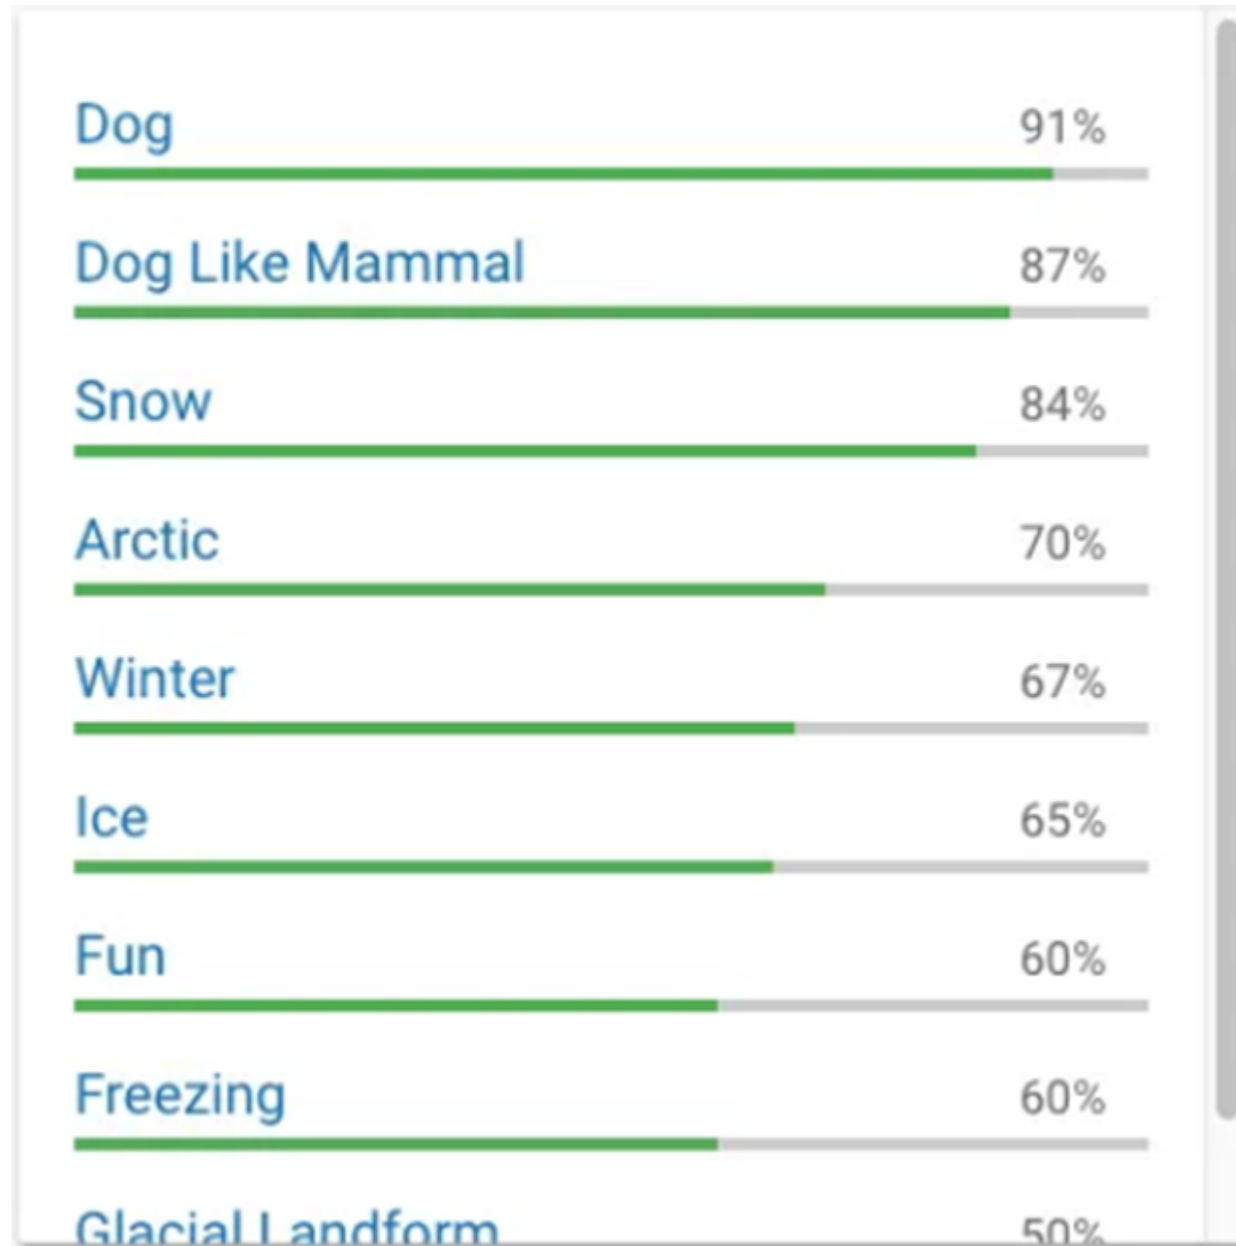




skier_adv.png

Human readers will easily identify the image as showing two men on skis. Google's Cloud Vision service reported being 91 percent certain it saw a dog. Other stunts have shown how to make stop signs invisible, or audio that sounds benign to humans but is transcribed by software as "OK Google browse to evil dot com."

LABSIX



Model #2: Adversarially Robust Streaming

- **Input:** Elements of an underlying data set S , which arrives sequentially and *adversarially*
- **Output:** Evaluation (or approximation) of a given function
- **Goal:** Use space *sublinear* in the size m of the input S



1

1



Model #2: Adversarially Robust Streaming

- **Input:** Elements of an underlying data set S , which arrives sequentially and *adversarially*
- **Output:** Evaluation (or approximation) of a given function
- **Goal:** Use space *sublinear* in the size m of the input S



10

1



Model #2: Adversarially Robust Streaming

- **Input:** Elements of an underlying data set S , which arrives sequentially and *adversarially*
- **Output:** Evaluation (or approximation) of a given function
- **Goal:** Use space *sublinear* in the size m of the input S



101

2



Model #2: Adversarially Robust Streaming

- **Input:** Elements of an underlying data set S , which arrives sequentially and *adversarially*
- **Output:** Evaluation (or approximation) of a given function
- **Goal:** Use space *sublinear* in the size m of the input S



1010

3



Model #2: Adversarially Robust Streaming

- **Input:** Elements of an underlying data set S , which arrives sequentially and *adversarially*
- **Output:** Evaluation (or approximation) of a given function
- **Goal:** Use space *sublinear* in the size m of the input S



1010

3



$(1 + \varepsilon)$ -Robust Algorithms [WZ21]

- $\tilde{\Theta}\left(\frac{1}{\varepsilon^2}\right)$ space-accuracy tradeoff for distinct elements, i.e., F_0
- $\tilde{\Theta}\left(\frac{1}{\varepsilon^2}\right)$ space-accuracy tradeoff for F_p with $p \in (0, 2]$
- $\tilde{\Theta}\left(\frac{1}{\varepsilon^2}\right)$ space-accuracy tradeoff for F_p with integer $p > 2$
- $\tilde{\Theta}\left(\frac{1}{\varepsilon^2}\right)$ space-accuracy tradeoff for L_2 -heavy hitters

“No losses* are necessary!”

15-859 ALGORITHMS FOR BIG DATA — Fall 2022

PROBLEM SET 3

Due: Tuesday, November 8, 11:59pm

Problem 2: F_2 -Difference Estimator in a Stream (25 points)

In this problem we consider insertion-only streams, meaning that we just see positive updates to an underlying vector $x \in \{0, 1, 2, \dots, M\}^n$ for some $M = \text{poly}(n)$. That is, no negative changes to coordinates of x are allowed. We will consider estimating $\|x\|_2^2$ up to a $(1 + \epsilon)$ -multiplicative factor.

Often x is *slowly-changing*, meaning that at some point in the stream $x = u$ and then at some later point in the stream $x = u + v$ for some $v \in \{0, 1, 2, \dots, M\}^n$, and we have $\|u + v\|_2^2 - \|u\|_2^2 \leq \gamma\|u\|_2^2$ and $\|v\|_2^2 \leq \gamma\|u\|_2^2$ for some $0 < \gamma < 1$. You would like to estimate $\|u + v\|_2^2$ using your previous estimate $\|u\|_2^2$ and using very little space. Show how to estimate $\|u + v\|_2^2$ up to a $(1 \pm \epsilon)$ -multiplicative factor with probability at least $9/10$, given a $(1 \pm \epsilon/2)$ -approximation to $\|u\|_2^2$ and a sketch which uses space $O(\gamma\epsilon^{-2} \log n)$ bits. You can assume that you initialize $x = 0$ and after processing some number of updates you will have $x = u$ at which point you are given a $1 \pm \epsilon/2$ approximation to $\|u\|_2^2$. Then you start processing the rest of the updates and finally end up setting $x = u + v$ and at the end of the stream you want to output a $1 \pm \epsilon$ approximation to $\|u + v\|_2^2$.

Note that if $\gamma = \Theta(1)$, then there is no improvement over just estimating $\|u + v\|_2^2$ directly using the sketch from class. However, if for example, $\gamma = \Theta(\epsilon)$, then the memory is only $O(\epsilon^{-1} \log n)$ bits, which is a significant savings.


Additional Work



- [BHMSSZ22] showed that “sampling-based” algorithms achieve adversarial robustness “for-free”
- [ABJSSWZ22] showed impossibility results for many natural problems for “white-box” adversaries
- [ABJSSWZ22] gave a sublinear-space algorithm for estimating F_0 when the white-box adversary has polynomial computation time
- **Open Questions:** 1) Is anything non-trivial possible for “black-box” adversaries that are allowed to delete elements? 2) Can lattice-based crypto be used for other algorithms robust to white-box adversaries?

Evolving Demands


- Sublinear-time or sublinear-space algorithms
- Incorporation of advice
- Security and privacy
- Robustness to noise or adversarial input
- **Ability to handle time-sensitive data**



“Facebook data policy retains user search histories for 6 months, the Apple differential privacy overview states that collected data is retained for 3 months, the Google data retention policy states that browser information may be stored for up to 9 months”

What is the data retention policy?

Data retention policies **concern what data should be stored or archived, where that should happen, and for exactly how long.** Once the retention time period for a particular data set expires, it can be deleted or moved as historical data to secondary or tertiary storage, depending on the requirements.



Do I need a data retention policy?

There are numerous benefits to establishing a solid data retention policy. Some of the more compelling benefits include: Automated compliance. With an established policy, organizations can ensure they comply with regulatory requirements mandating the retention of various types of data.

Model #3: Sliding Window Model

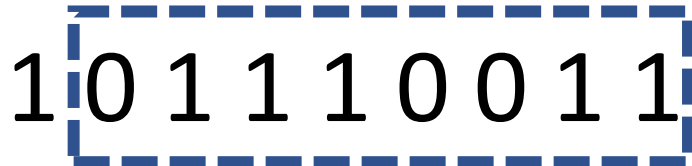
- **Input:** Elements of an underlying data set S , which arrives sequentially
- **Output:** Evaluation (or approximation) of a given function
- **Goal:** Use space *sublinear* in the size m of the input S
- **Sliding Window:** “Only the m most recent updates form the underlying data set S ”

1 0 1 1 1 0 0 1

Model #3: Sliding Window Model

- **Input:** Elements of an underlying data set S , which arrives sequentially
- **Output:** Evaluation (or approximation) of a given function
- **Goal:** Use space *sublinear* in the size m of the input S
- **Sliding Window:** “Only the m most recent updates form the underlying data set S ”

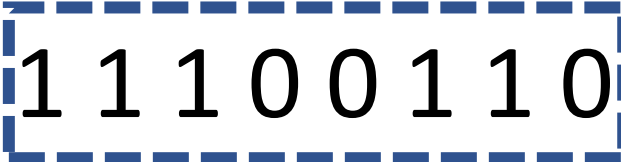
1 0 1 1 1 0 0 1 1



Model #3: Sliding Window Model

- **Input:** Elements of an underlying data set S , which arrives sequentially
- **Output:** Evaluation (or approximation) of a given function
- **Goal:** Use space *sublinear* in the size m of the input S
- **Sliding Window:** “Only the m most recent updates form the underlying data set S ”


1 0 1 1 1 0 0 1 1 0



Model #3: Sliding Window Model

- **Input:** Elements of an underlying data set S , which arrives sequentially
- **Output:** Evaluation (or approximation) of a given function
- **Goal:** Use space *sublinear* in the size m of the input S
- **Sliding Window:** “Only the m most recent updates form the underlying data set S ”
 - Emphasizes recent interactions, appropriate for time sensitive settings

1 0 1 1 1 0 0 1 1 0 1



$(1 + \varepsilon)$ -Approximation Sliding Window Algorithms [WZ21]

- Space $\tilde{\Theta}\left(\frac{1}{\varepsilon^2}\right)$ algorithm for F_p with $p \in (0, 2]$ [WZ21]

Problem	[BO07] Space	Our Result
L_p Estimation, $p \in (0, 1)$	$\tilde{O}\left(\frac{\log^3 n}{\varepsilon^3}\right)$	$\tilde{O}\left(\frac{\log^3 n}{\varepsilon^2}\right)$
L_p Estimation, $p \in (1, 2]$	$\tilde{O}\left(\frac{\log^3 n}{\varepsilon^{2+p}}\right)$	$\tilde{O}\left(\frac{\log^3 n}{\varepsilon^2}\right)$
L_p Estimation, integer $p > 2$	$\tilde{O}\left(\frac{n^{1-2/p}}{\varepsilon^{2+p}}\right)$	$\tilde{O}\left(\frac{n^{1-2/p}}{\varepsilon^2}\right)$
Entropy Estimation	$\tilde{O}\left(\frac{\log^5 n}{\varepsilon^4}\right)$	$\tilde{O}\left(\frac{\log^5 n}{\varepsilon^2}\right)$

“No losses* are necessary!”

$(1 + \varepsilon)$ -Approximation Sliding Window Algorithms [BGLWZ18]

- $\tilde{\Theta}\left(\frac{1}{\varepsilon^2}\right)$ space-accuracy tradeoff for F_0 , i.e., distinct elements
- $\tilde{\Theta}\left(\frac{1}{\varepsilon^2}\right)$ space-accuracy tradeoff for L_2 -heavy hitters
- We give nearly tight new upper and lower bounds for these problems!

RandNLA in the Sliding Window Model


- First algorithms for problems in numerical linear algebra in the sliding window model [BDMMUWZ20]
- $O(d^2)$ space randomized sliding window algorithm for spectral sparsification (space optimal, up to lower order terms)
- $O(dk)$ space randomized sliding window algorithm for low-rank approximation (space optimal, up to lower order terms)
- $O(d^3)$ space randomized sliding window algorithm for ℓ_1 subspace embedding

My Recent Work on Sublinear Algorithms

- Memory Bounds for the Experts Problem (STOC 2022)
- The White-Box Adversarial Data Stream Model (PODS 2022)
- Truly Perfect Samplers for Data Streams and Sliding Windows (PODS 2022)
- Noisy Boolean Hidden Matching with Applications (ITCS 2022)
- Adversarial Robustness of Streaming Algorithms through Importance Sampling (NeurIPS 2021)
- Tight Bounds for Adversarially Robust Streams and Sliding Windows via Difference Estimators (FOCS 2021)
- Separations for Estimating Large Frequency Moments on Data Streams (ICALP 2021)
- Non-Adaptive Adaptive Sampling on Turnstile Streams (STOC 2020)

Future Research Directions

- Adversarial robustness for data science/ML
- Limits of learning-augmented algorithms, beyond NP-hardness, better utility bounds for DP
- Data science with fairness considerations
- Evolving demands of the future!





Future Agenda

- **Department research:** interdisciplinary collaborations
- **Department activities:** social hour, reading group, lunch seminar
- **Regional activities:** Texas Area Theory Day, Texas Area Data Science Day

- **Teaching:** Randomized algorithms / mathematical toolkit (hashing, concentration inequalities, linear programming, convex optimization), “modern” data science (dimensionality reduction, clustering, sublinear algorithms, differential privacy)
- **Funding?:** NSF, NIH, DARPA, Apple, Google, Meta, Microsoft

Thank You!



Thank You!